

AI Infrastructure Earnings Call Review

Four quarters of architecture disclosures and unresolved questions

28 September 2026 | 35 companies | 140 quarterly events | 168 selected Q&A records

Executive summary

AI infrastructure decisions increasingly depend on the ability to deliver and operate a complete system. Across the selected calls, demand statements frequently provide less clarity about committed components, powered space, customer qualification and actual workload economics. Procurement should therefore connect each architectural choice to its delivery dependencies and measurable service outcome.

Plan the resources around the accelerator. CPU orchestration, host memory, storage and networking enter more of the discussion. A universal CPU-to-GPU ratio does not follow from these disclosures.

Evaluate inference by workload. Silicon specialization and deployment location depend on latency, throughput, model behavior, software and utilization.

Separate commitments from aspirations. Pipeline, orders, allocated components, energized capacity and accepted systems are different stages. Reported capex also depends on who funds the assets.

Require integrated acceptance. Power, cooling, networking and server delivery increasingly interact; interface ownership and commissioning matter alongside component performance.

How to read the evidence

This is a focused review of material Q&A, with at least one selected record for every company-quarter. It is not an exhaustive question census or a ranking of vendor transparency. Management answers remain attributed claims. The appendix preserves the questions, answer classifications, proposed follow-ups and transcript links.

Selected answer classifications

Classification	Records
Answered	75
Partially answered	91
Deferred	1
Unanswered	1

Architectural findings

1 Procure a commissioned system not a capex headline

Supply assurances refer to different stages: component allocation, equipment delivery, powered-shell completion, utility agreements and service acceptance. CoreWeave identifies shell execution in one delayed project; Equinix identifies power in dense facilities. Those statements concern different constraints and need not conflict.

Decision implication: Require a dated chain of commitments from wafers and memory through integration, energization and workload acceptance. Record the owner and recovery plan at each stage.

Sources: [CoreWeave FY2025 Q3](#); [Equinix FY2026 Q1](#); [HPE FY2026 Q1](#)

2 The CPU question is now part of AI capacity planning

Compute vendors increasingly describe orchestration and tool execution alongside accelerator work. Yet Dell says much conventional server growth still comes from established enterprise demand, and TSMC excludes CPUs from its AI accelerator measure. Revenue labels therefore do not identify a universal agentic CPU/GPU ratio.

Decision implication: Size CPU cores, DRAM, storage I/O and network demand from measured agent traces; separate refresh demand from new AI execution.

Sources: [NVIDIA FY2027 Q1](#); [Intel FY2026 Q1](#); [Dell FY2027 Q2](#); [TSMC FY2026 Q1](#)

3 Memory is a hierarchy and an allocation problem

HBM requirements are rising, while suppliers discuss conventional DRAM, flash, cache offload and stronger agreements. Additional HBM output competes for manufacturing resources with other products. Statements about samples, booked capacity and shipments represent different evidence levels.

Decision implication: Specify bandwidth, capacity, latency, endurance and power by tier. Ask for qualified delivery quantities and contract terms, not only a claim that annual supply is sold out.

Sources: [Micron FY2026 Q1](#); [Micron FY2026 Q2](#); [SK hynix FY2026 Q2](#); [Samsung FY2025 Q3](#)

Deployment and procurement implications

4 Inference does not have one winning chip or location

Amazon and Microsoft retain multiple silicon choices while investing internally. Oracle emphasizes hardware execution latency; Cisco emphasizes local real-time use cases. These positions reflect different workloads rather than a settled central-versus-edge outcome.

Decision implication: Compare complete request paths: prefill, decode, tools, retrieval, network transit and service-level targets. Include software migration, utilization and fallback costs.

Sources: [Amazon FY2025 Q3](#); [Microsoft FY2026 Q2](#); [Oracle FY2026 Q3](#); [Cisco FY2026 Q1](#)

5 Architecture boundaries are moving toward integrated delivery

Networking, cooling and power suppliers describe increasing interdependencies. Vertiv explicitly calls the procurement shift gradual and partial. Eaton describes a broader capability set for DC infrastructure, while Schneider expects staged adoption.

Decision implication: Make interface ownership, qualification, protection, maintainability and acceptance criteria part of the system contract. Do not infer universal adoption from a reference design.

Sources: [Vertiv FY2026 Q1](#); [Eaton FY2026 Q2](#); [Schneider Electric FY2026 Q1](#)

6 Growth statistics need a physical and contractual denominator

Lenovo distinguishes opportunities from orders. Oracle and Meta describe financing structures that change reported capex without eliminating infrastructure investment. Equipment suppliers distinguish manufacturing volume from process intensity and upgrades.

Decision implication: Compare committed versus prospective demand; units versus price; energized MW versus land runway; and all-party investment versus issuer capex. Avoid adding supplier revenues across the same stack.

Sources: [Lenovo FY2027 Q1](#); [Oracle FY2027 Q1](#); [Meta FY2025 Q3](#); [ASML FY2026 Q2](#)

Priority follow up questions

These questions are proposed diligence actions, ordered by architectural consequence. They are not claims that every analyst asked them.

Area	Question and evidence required
1 Capacity delivery	Which quantities are contractually allocated, installed, energized and accepted in each quarter? A dated milestone table, cancellation rights, responsible party and recovery alternatives.
2 Fleet economics	What is cost per successful task at a specified latency and quality target? Matched workload, utilization, software, power, networking and amortization assumptions.
3 CPU and memory	What does an agent trace consume outside the accelerator? CPU time, host memory, cache traffic, storage I/O and concurrency distributions.
4 HBM and packaging	Which customer-qualified products can ship together, and when? Qualification, yield, package availability and accepted shipment quantities.
5 Inference placement	Which latency component justifies central, regional or on-device execution? Measured network, prefill, decode, retrieval and tool latencies, plus privacy requirements.
6 Power and cooling	Who owns the complete electrical and thermal design across a generation change? Protection studies, operating envelopes, redundancy, service and integration acceptance.
7 Demand quality	How much pipeline becomes non-cancellable orders and operating revenue? Cohort conversion, concentration, financing and deployment dependencies.
8 Existing asset value	What workloads remain economic on older hardware? Utilization, maintenance, impairment and workload reassignment evidence.

Coverage and method

The review covers seven compute suppliers, three memory suppliers, TSMC, nine cloud and platform companies, four server suppliers, two networking suppliers, three power and cooling companies, two colocation operators and four semiconductor equipment suppliers. Each has four quarterly slots. Fiscal labels are retained, so the periods are not identical calendar windows.

Source handling

Company-hosted transcripts were used where available, supplemented by transcript mirrors of management testimony. Mislabelled or unsuitable entries were replaced, including ASML media or prerecorded calls, a Lam Research entry containing another company's call, and an inconsistent Applied Materials entry. Recordings were not systematically replayed. Translated transcripts and uncertain speaker labels remain limitations.

Selection and classification

The 168 records are selected question-level probes, not a complete set of normalized atomic asks. Related subparts sometimes remain grouped. Compute suppliers have more selected probes than most other companies. Counts describe this ledger only and cannot measure audience interest, investor sentiment or relative transparency.

Answered means the selected substantive ask is directly addressed. **Partially answered** means a requested element remains unresolved. **Deferred** records explicit postponement. **Unanswered** includes a refusal to provide the requested disclosure. No selected record was classified as deflected. A proposed validation step does not by itself make a direct answer partial.

Scope boundaries

The universe is broad but does not include every supplier. Private model labs and infrastructure vendors lack comparable public quarterly Q&A. Additional ODMs, specialist interconnect suppliers, storage-only vendors, OSATs and utilities are outside this edition. Schneider's Q1 and Q3 calls are revenue updates, alongside half-year and full-year results. Qualcomm is included for its emerging infrastructure role; some selected discussion concerns adjacent device markets.

Using the appendix

Each company starts on a new page. Quarter links open the source transcript. The record identifier matches the interactive ledger. Questions and answers are paraphrases; the questioner helps locate the original exchange. Status applies at the call date. The follow-up column combines unresolved parts of the question and additional architectural validation proposed by this review.

Company index

Company	Layer	Appendix page
NVIDIA	Compute	7
AMD	Compute	9
Intel	Compute	11
Arm	Compute	13
Broadcom	Compute	15
Marvell	Compute	17
Qualcomm	Compute	19
Micron	Memory	21
SK hynix	Memory	22
Samsung	Memory	23
TSMC	Foundry	24
Microsoft	Cloud	25
Amazon	Cloud	26
Alphabet	Cloud	27
Meta	Cloud	28
Oracle	Cloud	29
Alibaba	Cloud	30
Tencent	Cloud	31
Baidu	Cloud	32
CoreWeave	Cloud	33
Dell	Systems	34
HPE	Systems	35
Supermicro	Systems	36
Lenovo	Systems	37
Arista	Network	38
Cisco	Network	39
Vertiv	Power cooling	40
Eaton	Power cooling	41
Schneider Electric	Power cooling	42
Equinix	Colocation	43
Digital Realty	Colocation	44
ASML	Equipment	45
Applied Materials	Equipment	46
Lam Research	Equipment	47
KLA	Equipment	48

NVIDIA

Compute | Four quarter comparison

CPU orchestration and inference specialization broaden the platform story; requests to rank supply bottlenecks still lack a clear hierarchy.

FY2026 Q3 | November 19, 2025

QA001 | Supply | Partially answered

Question Which constraint could limit growth most?

Management Management identified upstream supply and downstream sites, power and financing; no bottleneck ranking.

Follow up Relative constraint severity.

Questioner: Timothy Arcuri

QA002 | Power economics | Partially answered

Question What power and content assumptions underpin growth?

Management Management tied rising platform content to generational efficiency and fixed site power.

Follow up Assumptions behind the aggregate forecast.

Questioner: Vivek Arya

FY2026 Q4 | February 25, 2026

QA003 | CPU | Answered

Question Why sell Vera separately?

Management Management described bandwidth-rich CPUs for preprocessing, tools and post-training.

Follow up Independent workload comparison remains needed.

Questioner: Aaron Rakers

QA004 | Inference | Answered

Question Which applications drive the next demand wave?

Management Management emphasized agentic coding, followed by physical AI.

Follow up Application adoption and capacity forecasts remain claims.

Questioner: James Schneider

NVIDIA continued

Later two quarters

FY2027 Q1 | May 20, 2026

QA005 | CPU | Answered

Question Do agentic CPUs replace GPU demand?

Management Management described CPU harnesses and tools generating additional GPU inference; standalone demand was distinguished.

Follow up Fleet-specific CPU/GPU provisioning ratio.

Questioner: Vivek Arya

QA006 | Platform | Answered

Question Why separate hyperscale from other customers?

Management Management emphasized diverse customers needing integrated infrastructure.

Follow up Comparable segment deployment economics.

Questioner: Joseph Moore

FY2027 Q2 | August 26, 2026

QA007 | Supply | Partially answered

Question Rank the most acute deployment constraints.

Management Management described broad supply-chain pressure without ranking bottlenecks.

Follow up Constraint ranking and relief dates.

Questioner: James Schneider

QA008 | Inference | Partially answered

Question How do agentic workloads and Groq affect inference position?

Management Management promoted platform flexibility; Groq targets high interactivity with throughput/cost tradeoffs.

Follow up Quantified inference share.

Questioner: Christopher Muse

Transcript links above provide the full exchange. Management claims are not independent performance validation.

AMD

Compute | Four quarter comparison

Agentic CPU demand and accelerator ramps strengthen the narrative, while detailed workload comparisons and secured capacity remain open.

FY2025 Q3 | November 4, 2025

QA009 | Supply | Partially answered

Question Can components and power support deployment plans?

Management Management described coordinated two-year planning with customers and suppliers.

Follow up Project-specific delivery assurances.

Questioner: Thomas O'Malley

QA010 | CPU | Answered

Question Is server demand durable beyond seasonality?

Management Management reported broadening hyperscaler CPU plans into 2026.

Follow up Attribution to agentic versus conventional workloads.

Questioner: Joshua Buchalter

FY2025 Q4 | February 3, 2026

QA011 | Inference | Partially answered

Question How does Instinct compete with specialized low-latency silicon?

Management Management pointed to chiplet flexibility and workload optimization across inference stages.

Follow up Measured latency and cost comparison.

Questioner: Joshua Buchalter

QA012 | Supply | Partially answered

Question What China assumptions and data-center drivers support growth?

Management Management limited modeled China sales and cited EPYC and accelerator ramps.

Follow up Complete growth bridge.

Questioner: Vivek Arya

AMD continued

Later two quarters

FY2026 Q1 | May 5, 2026

QA013 | Supply | Partially answered

Question Can supply support CPU and GPU ramp cadence?

Management Management acknowledged tightness but asserted support for planned growth and coordinated power visibility.

Follow up Committed deployment capacity beyond guidance.

Questioner: Thomas O'Malley

QA014 | CPU | Partially answered

Question Why expand the CPU market outlook despite competition?

Management Management attributed faster demand to agent orchestration and data processing.

Follow up Workload-level basis for market and share forecasts.

Questioner: Joshua Buchalter

FY2026 Q2 | August 4, 2026

QA015 | CPU | Partially answered

Question When will accelerator revenue overtake CPUs?

Management Management emphasized growth in both and the larger accelerator opportunity.

Follow up Crossover timing.

Questioner: Thomas O'Malley

QA016 | Supply | Partially answered

Question What drives the CPU/GPU growth cadence?

Management Management linked added supply to CPU growth and a larger fourth-quarter Helios ramp.

Follow up Exact CPU/GPU contribution split.

Questioner: Timothy Arcuri

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Intel

Compute | Four quarter comparison

Foundry discussions move toward readiness and investment, but customer engagement, qualified production and sustained volume remain distinct milestones.

FY2025 Q3 | October 23, 2025

QA017 | Platform | Partially answered

Question How does NVIDIA collaboration fit inference plans?

Management Management discussed NVLink connectivity and future inference products.

Follow up Detailed product integration.

Questioner: Aaron Rakers

QA018 | Foundry | Partially answered

Question How can customer commitments precede capacity investment?

Management Management described engagement and reuse of existing footprint and equipment.

Follow up Binding commitments and qualified output.

Questioner: Joseph Moore

FY2025 Q4 | January 22, 2026

QA019 | CPU | Partially answered

Question Does unmet server demand migrate to competitors?

Management Management framed demand as an x86 refresh and expected supply improvement.

Follow up Market sizing and share migration.

Questioner: Vivek Arya

QA020 | Foundry | Answered

Question When will Intel loosen capacity spending?

Management Management emphasized yield/throughput gains and customer-backed 14A capacity investment.

Follow up Contract details and achieved yields.

Questioner: Ross Seymore

Intel continued

Later two quarters

FY2026 Q1 | April 23, 2026

QA021 | CPU | Answered

Question How should agentic CPU demand be modeled?

Management Management offered illustrative GPU/CPU ratios declining from training to agentic workloads.

Follow up Ratios are not a universal fleet sizing rule.

Questioner: Joshua Buchalter

QA022 | CPU | Partially answered

Question How does Intel compete with AMD and Arm?

Management Management cited architecture optimization and foundry opportunities without a quantified share answer.

Follow up Comparable performance and market share.

Questioner: Vivek Arya

FY2026 Q2 | July 23, 2026

QA023 | Foundry | Partially answered

Question Does higher capex imply external foundry orders?

Management Management described broad investment, packaging and long-term agreements.

Follow up Customer-specific commitments and allocation.

Questioner: Benjamin Reitzes

QA024 | Foundry | Partially answered

Question What supports external-foundry confidence and spending?

Management Management cited process readiness and PDK milestones; engagement remained distinct from production.

Follow up Binding volume and capex split.

Questioner: Vivek Arya

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Arm

Compute | Four quarter comparison

CPU positioning broadens into system-level supply questions; architectural efficiency claims still need workload-level validation.

FY2026 Q2 | November 5, 2025

QA025 | Inference | Answered

Question How does inference expansion affect Arm?

Management Management described diverse inference architectures extending from cloud to edge.

Follow up Deployment mix remains unquantified.

Questioner: Sreekrishnan Sankarnarayanan

QA026 | Power economics | Answered

Question How is Arm positioned for new AI capacity?

Management Management used power constraints to argue for Arm efficiency and adoption.

Follow up Efficiency assertions need workload normalization.

Questioner: Sebastien Cyrus Naji

FY2026 Q3 | February 4, 2026

QA027 | Memory | Partially answered

Question What do new memory technologies mean for Arm?

Management Management described ongoing memory and power-efficiency research without a product schedule.

Follow up Specific technology and timing.

Questioner: Timm Schulze-Melander

QA028 | CPU | Answered

Question How does inference change the CPU role?

Management Management described more CPU deployment for efficient, low-latency work.

Follow up Workload-specific CPU demand.

Questioner: Joe Quatrochi

Arm continued

Later two quarters

FY2026 Q4 | May 6, 2026

QA029 | Supply | Partially answered

Question Can supply catch up with AGI CPU demand?

Management Management distinguished demand from secured wafer, memory, packaging and test capacity.

Follow up Incremental supply commitment.

Questioner: Andrew Gardiner

QA030 | CPU | Partially answered

Question When might CPU/GPU ratios pass parity?

Management Management emphasized separate CPU capacity beyond GPU head nodes.

Follow up Comparable counting units and timing.

Questioner: Sreekrishnan Sankarnarayanan

FY2027 Q1 | July 29, 2026

QA031 | CPU | Partially answered

Question How will core counts and silicon reporting evolve?

Management Management favored more cores but withheld detailed roadmap commitments.

Follow up Specific roadmap and reporting details.

Questioner: Vijay Rakesh

QA032 | Supply | Answered

Question What improved confidence in supplying AGI CPUs?

Management Management cited progress across wafers, substrates, memory and testing.

Follow up Supplier-specific allocations remain undisclosed.

Questioner: Joseph Quatrochi

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Broadcom

Compute | Four quarter comparison

Custom accelerators increasingly connect to system delivery, packaging and site readiness; specialization is customer- and workload-dependent.

FY2025 Q4 | December 11, 2025

QA033 | Supply | Partially answered

Question Are key materials secured for the backlog?

Management Management described additional capacity work and advanced packaging investment.

Follow up Component-by-component supply coverage.

Questioner: Harlan Sur

QA034 | Platform | Answered

Question Are new customer deliveries chips or systems?

Management Management explicitly described responsibility for system sales.

Follow up Customer acceptance criteria.

Questioner: Blayne Curtis

FY2026 Q1 | March 4, 2026

QA035 | Custom silicon | Partially answered

Question Could customer-owned designs displace Broadcom?

Management Management defended integrated silicon, SerDes, packaging and networking expertise.

Follow up Displacement magnitude.

Questioner: Harlan Sur

QA054 | Custom silicon | Answered

Question How does prefill/decode specialization change custom silicon competition?

Management Management described specialization by customer, model and workload across its customer roadmaps.

Follow up Compare end-to-end performance and cost; initial question 4 was clarified in question 5.

Questioner: Christopher Muse

Broadcom continued

Later two quarters

FY2026 Q2 | Jun 3, 2026

QA036 | Supply | Answered

Question Can Broadcom obtain more wafers and HBM?

Management Management said requirements for 2026-27 were secured and later years under planning.

Follow up Incremental upside capacity.

Questioner: Timothy Arcuri

QA037 | CPU | Answered

Question Does agentic AI change infrastructure-software demand?

Management Management linked growing CPU core deployment to VMware demand.

Follow up Workload attribution versus licensing effects.

Questioner: Atif Malik

FY2026 Q3 | September 2, 2026

QA038 | Custom silicon | Partially answered

Question Why do frontier labs seek direct compute capacity?

Management Management described lab economics and preference for first-party optimized infrastructure.

Follow up Comparable economics and customer mix.

Questioner: Vivek Arya

QA039 | Supply | Answered

Question Which constraints govern XPU deployments?

Management Management identified sites, power, HBM and customer-procured system memory.

Follow up Project-level bottleneck ranking.

Questioner: James Schneider

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Marvell

Compute | Four quarter comparison

The opportunity expands beyond the main accelerator into optics, interconnect and attached memory functions; content mix remains incompletely disclosed.

FY2026 Q3 | December 2, 2025

QA040 | Network | Answered

Question Does Celestial's outlook include memory products?

Management Management included the whole business but expected photonic chiplets to lead.

Follow up Revenue timing remains a forecast.

Questioner: Tore Svanberg

QA041 | Network | Answered

Question Why use cloud capex to model optics?

Management Management called capex a broad proxy; AI optics could grow faster.

Follow up Customer-level optical content.

Questioner: Vivek Arya

FY2026 Q4 | March 5, 2026

QA042 | Custom silicon | Partially answered

Question What drives custom compute and attached content?

Management Management described participation beyond the main XPU and new customer programs.

Follow up Program-specific revenue allocation.

Questioner: Harlan Sur

QA043 | Network | Answered

Question Can optics outgrow hyperscaler capex?

Management Management linked growth to accelerator generations and connectivity demand.

Follow up Port and bandwidth mix.

Questioner: Aaron Rakers

Marvell continued

Later two quarters

FY2027 Q1 | May 27, 2026

QA044 | Inference | Partially answered

Question Are SRAM offload accelerators a design opportunity?

Management Management emphasized SRAM IP and custom-design capabilities.

Follow up Specific wins and performance.

Questioner: Harlan Sur

QA045 | Custom silicon | Answered

Question What does the custom-silicon outlook include?

Management Management distinguished the overall custom market from XPU-only content.

Follow up Customer conversion of forecasts.

Questioner: Vivek Arya

FY2027 Q2 | August 27, 2026

QA046 | Memory | Answered

Question How do memory constraints change custom designs?

Management Management cited CXL expansion and inference-specific acceleration.

Follow up Measured system benefit.

Questioner: Christopher Muse

QA047 | Custom silicon | Partially answered

Question What products sit within the expanded engagement?

Management Management named inference, NIC, storage, memory-interface and near-memory compute opportunities.

Follow up Contribution by product.

Questioner: Thomas O'Malley

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Qualcomm

Compute | Four quarter comparison

Roadmap discussion progresses to reported purchase orders and wafer starts; broader adoption still requires demonstrated silicon and customer qualification.

FY2025 Q4 | November 5, 2025

QA048 | Inference | Answered

Question Why depart from conventional GPU/HBM architectures?

Management Management prioritized inference density, cost and energy through alternative designs.

Follow up Validated performance and software readiness.

Questioner: Christopher Caso

QA049 | Inference | Partially answered

Question What are deployment economics and ramp timing?

Management Management discussed an earlier revenue ramp but limited comparative economics.

Follow up Normalized system economics.

Questioner: Samik Chatterjee

FY2026 Q1 | February 4, 2026

QA050 | Memory | Answered

Question How can Qualcomm manage memory scarcity?

Management Management cited broad qualification while noting customers procure most handset memory.

Follow up Handset answer; indirect infrastructure evidence.

Questioner: C. J. Muse

QA051 | Inference | Partially answered

Question How are data-center programs progressing?

Management Management described disaggregated decode, CPU options and a new memory architecture.

Follow up Customer delivery milestones.

Questioner: Samik Chatterjee

Qualcomm continued

Later two quarters

FY2026 Q2 | April 29, 2026

QA052 | Custom silicon | Deferred

Question What is the hyperscaler engagement's scope?

Management Management withheld details until the June investor event.

Follow up Scope and product mix at this call.

Questioner: Joshua Buchalter

QA055 | Inference | Partially answered

Question How will Qualcomm compete in data-center inference?

Management Management outlined disaggregated inference, CPUs and custom silicon; further roadmap disclosure was scheduled later.

Follow up Product-level comparison and delivery milestones.

Questioner: Samik Chatterjee

FY2026 Q3 | July 29, 2026

QA053 | Supply | Answered

Question Are wafer and assembly constraints limiting plans?

Management Management reported broad tightness but asserted supply support for its plans.

Follow up Headroom beyond current plans.

Questioner: Joseph Moore

QA056 | Supply | Answered

Question What supports the hyperscaler revenue ramp?

Management Management cited two engagements, purchase orders and wafer starts, with revenue beginning in the December quarter.

Follow up Track shipment acceptance and subsequent ramp.

Questioner: Joshua Buchalter

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Micron

Memory | Four quarter comparison

The discussion broadens from HBM performance to stronger supply agreements, alternative memory tiers and persistent limits on supply growth.

FY2025 Q4 | September 23, 2025

QA057 | Memory | Answered

Question How will HBM meet customer bandwidth requirements above standard specifications?

Management Management attributed its capabilities to DRAM architecture, circuit design and the base die.

Follow up Validate customer-qualified bandwidth and power; transcript units are unreliable.

Questioner: Harlan Sur

FY2026 Q1 | December 17, 2025

QA058 | Supply | Partially answered

Question How firm and broad are memory supply agreements?

Management Management described multiyear DRAM and NAND negotiations with stronger commitments, withholding detailed terms.

Follow up Binding quantities, price mechanisms and cancellation rights.

Questioner: Timothy Arcuri

FY2026 Q2 | March 18, 2026

QA059 | Memory | Answered

Question What is the opportunity and limitation of high-bandwidth flash?

Management Management acknowledged capacity benefits but identified write-speed, power and retention limitations.

Follow up Benchmark workload suitability against other memory tiers.

Questioner: Harlan L. Sur

FY2026 Q3 | June 24, 2026

QA060 | Supply | Partially answered

Question How much worse will memory undersupply be next year?

Management Management expected persistent tightness and explained slow capacity additions without sizing the future gap.

Follow up Demand-supply gap by memory type and delivery quarter.

Questioner: Krish Sankar

Transcript links above provide the full exchange. Management claims are not independent performance validation.

SK hynix

Memory | Four quarter comparison

Higher HBM requirements, stronger agreements and balanced allocation evolve into a more explicit, workload-specific SSD hierarchy.

FY2025 Q3 | Oct 29, 2025

QA154 | Memory | Answered

Question Can HBM meet above-standard customer requirements?

Management Management said enhanced samples and production address higher bandwidth needs and expected requirements to keep rising.

Follow up Verify customer qualification and bandwidth/power together.

Questioner: Kim Ki-tae / operator names Jong-wook Lee; attribution uncertain

FY2025 Q4 | Jan 29, 2026

QA155 | Supply | Partially answered

Question How do new memory supply agreements differ from older ones?

Management Management described stronger mutual commitments than historically flexible volume agreements, withholding contract specifics.

Follow up Binding term, volume, price and cancellation provisions.

Questioner: Hyung Kyun Ryu, Daishin Securities

FY2026 Q1 | Apr 23, 2026

QA156 | Memory | Partially answered

Question What is HBM4 qualification status, shipment timing and share?

Management Management emphasized customer-aligned ramp readiness and overall competitiveness without the requested detailed schedule or share.

Follow up Customer acceptance, mass shipments and product mix.

Questioner: Ricky Seo, HSBC

FY2026 Q2 | Jul 29, 2026

QA157 | Memory | Answered

Question How will SSD products address inference and cache offload?

Management Management described a workload-specific hierarchy of high-capacity QLC, performance TLC and SLC-mode products.

Follow up Measured cache-offload latency, endurance and total cost.

Questioner: Hyung-keun Ryu, Daishin Securities

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Samsung

Memory | Four quarter comparison

Earlier customer-qualification confidentiality gives way to ramp and allocation discussion; commercial acceptance still needs customer-specific confirmation.

FY2025 Q3 | October 30, 2025

QA158 | Memory | Partially answered

Question Has customer HBM qualification completed and what is the roadmap?

Management Management withheld customer-specific qualification under confidentiality, while describing HBM3E sales and HBM4 samples.

Follow up Customer-specific acceptance is not established by shipment growth alone.

Questioner: Dongwon Kim, KB Securities

FY2025 Q4 | January 29, 2026

QA151 | Foundry | Partially answered

Question What is the status of advanced nodes and turnkey memory/packaging?

Management Management described process targets, customer collaboration and integrated capabilities.

Follow up Customer qualification, yields and production volumes.

Questioner: Nicolas Gaudois

FY2026 Q1 | April 29, 2026

QA152 | Memory | Answered

Question How is HBM4 ramping?

Management Management described commercial shipments, booked capacity and a larger second-half contribution.

Follow up Independent customer acceptance and realized ramp.

Questioner: S. K. Kim

FY2026 Q2 | July 29, 2026

QA153 | Memory | Answered

Question How will Samsung balance HBM and conventional DRAM?

Management Management maintained a balanced portfolio objective while expanding HBM and citing future customer agreements.

Follow up Wafer allocation and opportunity cost by product.

Questioner: Ricky Seo

Transcript links above provide the full exchange. Management claims are not independent performance validation.

TSMC

Foundry | Four quarter comparison

Questions shift from translating AI spending into wafers toward shared leading-node allocation among CPUs, GPUs and custom accelerators.

FY2025 Q3 | October 16, 2025

QA061 | Capacity | Unanswered

Question How much wafer demand corresponds to a gigawatt of AI capacity?

Management Management declined a conversion because project configurations differ.

Follow up Project-specific chip mix and wafer content; no universal GW conversion.

Questioner: Sreekrishnan Sankarnarayanan

FY2025 Q4 | January 15, 2026

QA062 | Capacity | Partially answered

Question Will new capital spending relieve near-term shortages, and is engineering labor limiting?

Management Management placed major new-fab supply later, emphasizing near-term productivity; labor shortage was not quantified.

Follow up Separate productivity gains, new capacity and staffing limits.

Questioner: Brett Simpson

FY2026 Q1 | April 16, 2026

QA063 | CPU | Answered

Question Will CPUs enter TSMC's AI revenue definition?

Management Management said CPUs remain excluded because it cannot reliably identify their destination; historical AI figures were not restated.

Follow up Preserve the narrower accelerator definition in comparisons.

Questioner: Yu Shi

FY2026 Q2 | July 16, 2026

QA064 | Capacity | Partially answered

Question What is the relative growth and capacity allocation for CPUs and accelerators?

Management Management said they share leading-edge technologies and allocation is coordinated with customers; no specific mix was provided.

Follow up Node-level CPU/GPU/XPU allocations.

Questioner: Laura Chen

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Microsoft

Cloud | Four quarter comparison

Fleet flexibility persists alongside internal silicon; the unresolved issue is converting spending and efficiency into usable capacity and cohort returns.

FY2026 Q1 | October 29, 2025

QA077 | Platform | Answered

Question How does fleet strategy support durable infrastructure economics?

Management Management emphasized geographically fungible capacity across research, internal products and external customers.

Follow up Measure actual portability and stranded-capacity exposure.

Questioner: Kasthuri Rangan

FY2026 Q2 | January 28, 2026

QA078 | Custom silicon | Partially answered

Question How does internal silicon affect inference economics?

Management Management described model-system co-design while retaining external suppliers for fleet-level cost flexibility.

Follow up Comparable workload costs and margin contribution.

Questioner: Mark Murphy

FY2026 Q3 | April 29, 2026

QA079 | Capacity | Partially answered

Question Can physical constraints support the spending ramp and allocation needs?

Management Management expressed procurement confidence but expected tension between internal and Azure demand to persist.

Follow up Accepted capacity by component and customer category.

Questioner: Karl Keirstead

FY2026 Q4 | July 29, 2026

QA080 | Economics | Partially answered

Question How have infrastructure returns and efficiency opportunities changed?

Management Management cited silicon, model choice, token efficiency and fleet operations as levers without a quantitative return comparison.

Follow up Realized returns on successive capacity cohorts.

Questioner: Gabriela Borges

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Amazon

Cloud | Four quarter comparison

Trainium adoption depends on software and volume execution; later questions emphasize procurement, inflation and backlog-to-capacity conversion.

FY2025 Q3 | October 30, 2025

QA065 | Custom silicon | Answered

Question What must Trainium overcome to broaden adoption?

Management Management identified volume delivery, software ecosystem improvement and customer proof points, alongside continued NVIDIA purchases.

Follow up Independent workload and migration-cost comparison.

Questioner: Brian Nowak

FY2025 Q4 | February 5, 2026

QA066 | Custom silicon | Partially answered

Question How is Rainier progressing and how should its chip count be understood?

Management Management described Anthropic training and additional Trainium use outside Rainier; a complete count reconciliation was not supplied.

Follow up Separate cluster count, other deployments and accepted capacity.

Questioner: Douglas Anmuth

FY2026 Q1 | April 29, 2026

QA067 | Memory | Partially answered

Question How will memory and storage inflation affect investment?

Management Management cited early supplier procurement and confidence in supply, but did not isolate the capital-spending effect.

Follow up Price-versus-volume bridge and contract protection.

Questioner: Shweta Khajuria

FY2026 Q2 | July 30, 2026

QA068 | Capacity | Partially answered

Question How does larger backlog change future capacity and input-cost exposure?

Management Management said known backlog enters planning and discussed contract versus on-demand demand; later capacity was not quantified.

Follow up Backlog-to-commissioned-capacity schedule and cost pass-through.

Questioner: Ken Gawrelski

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Alphabet

Cloud | Four quarter comparison

Integrated models and silicon support the strategy, while persistent constraints make compute allocation and investment lead times central.

FY2025 Q3 | October 29, 2025

QA069 | Custom silicon | Partially answered

Question How does custom silicon improve efficiency and external monetization?

Management Management linked its integrated stack and TPU/GPU offerings to demand and margin opportunities.

Follow up Incremental economics attributable specifically to silicon.

Questioner: Eric Sheridan

FY2025 Q4 | February 4, 2026

QA070 | Capacity | Partially answered

Question When will investment close the compute supply-demand gap?

Management Management expected constraints to persist, emphasizing supply-chain lead times and prior investment.

Follow up Capacity delivery schedule and gap closure date.

Questioner: Eric Sheridan

FY2026 Q1 | April 29, 2026

QA071 | Capacity | Partially answered

Question Where will scarce next-generation compute be applied?

Management Management emphasized Search, Gemini and agentic experiences, plus cloud and selected TPU hardware sales.

Follow up Allocation weights and return thresholds.

Questioner: Brian Nowak

FY2026 Q2 | July 22, 2026

QA072 | Capacity | Partially answered

Question How are returns and future spending assessed under supply constraints?

Management Management described broad opportunities and continuing demand above capacity without a quantified return bridge.

Follow up Incremental usable capacity and realized return metrics.

Questioner: Brian Nowak

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Meta

Cloud | Four quarter comparison

Compute supports both existing recommendations and new agents; financing form and buy-versus-sell capacity choices complicate simple capex comparisons.

FY2025 Q3 | October 29, 2025

QA073 | Financing | Answered

Question How does the data-center joint venture affect reported investment?

Management Management distinguished construction spending outside reported capex from Meta's ownership contributions and described flexibility.

Follow up Compare economic commitments rather than capex alone.

Questioner: Youssef Squali

FY2025 Q4 | January 28, 2026

QA074 | Inference | Partially answered

Question How does more compute translate into advertising performance?

Management Management described transferring knowledge from larger models to smaller runtime models and continuing capacity constraints.

Follow up Marginal performance and inference cost versus compute.

Questioner: Eric James Sheridan

FY2026 Q1 | April 29, 2026

QA075 | Inference | Partially answered

Question What is the roadmap for consumer and business agents?

Management Management distinguished nearer-term engagement and advertising work from longer-term interacting agents.

Follow up Deployment milestones and infrastructure requirements.

Questioner: Eric Sheridan

FY2026 Q2 | July 29, 2026

QA076 | Capacity | Partially answered

Question Why consider selling compute while buying external capacity?

Management Management described a portfolio of near-term monetization and longer-term internal intelligence investments.

Follow up Timing, chip and workload differences were not disentangled.

Questioner: Douglas Anmuth

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Oracle

Cloud | Four quarter comparison

Inference placement and contract design join financing in the discussion; customer or partner hardware can decouple growth from Oracle-funded capex.

FY2026 Q2 | December 10, 2025

QA093 | Financing | Partially answered

Question How much funding is required for AI capacity?

Management Management described different hardware ownership and supplier rental models rather than a single funding requirement.

Follow up Total economic commitment across financing models.

Questioner: Brad Zelnick

FY2026 Q3 | March 10, 2026

QA094 | Inference | Answered

Question Will inference require data centers closer to users?

Management Management argued that hardware execution currently dominates latency and emphasized specialized inference architectures.

Follow up Measure network versus compute latency for each workload and geography.

Questioner: Mark Murphy

FY2026 Q4 | June 10, 2026

QA095 | Supply | Answered

Question How are long-term contracts protected from component inflation?

Management Management described fixed prices when costs are secured and pass-through mechanisms when they are not.

Follow up Confirm coverage and exceptions in actual contracts.

Questioner: John DiFucci

FY2027 Q1 | September 10, 2026

QA096 | Financing | Partially answered

Question Can spending slow without limiting growth, and when does cash flow turn positive?

Management Management separated Oracle-funded capex from partner/customer capital and declined a specific cash-flow date.

Follow up Total capital deployed and cash-flow timing.

Questioner: Brad Zelnick

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Alibaba

Cloud | Four quarter comparison

Management describes priority for enterprise monetization and CPU procurement, then offers a payback model whose assumptions require validation.

FY2026 Q2 | November 25, 2025

QA081 | Capacity | Answered

Question How is scarce compute allocated across models, services and GPU rental?

Management Management prioritized higher-value model and inference services over simple GPU rental and described broad component shortages.

Follow up Quantify allocation and realized returns by service.

Questioner: Ellie Jiang

FY2026 Q3 | March 19, 2026

QA082 | Platform | Partially answered

Question How does Token Hub change coordination and priorities?

Management Management emphasized tighter application-model integration for agents rather than a numerical priority hierarchy.

Follow up Operational allocation rules and measurable integration benefits.

Questioner: Robin Zhu

FY2026 Q4 | May 13, 2026

QA083 | Capacity | Answered

Question How are resources split between enterprise and consumer AI?

Management Management said most infrastructure currently supports business customers because willingness to pay is clearer.

Follow up Track allocation as consumer demand develops.

Questioner: Ellie Jiang

FY2027 Q1 | August 20, 2026

QA084 | Economics | Partially answered

Question What drove spending and what return supports it?

Management Management cited delivery timing, CPU demand and component inflation, offering a payback framework based on current margins.

Follow up Auditable service-level allocation and payback assumptions; management's framework is not validated return.

Questioner: Alicis a Yap

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Tencent

Cloud | Four quarter comparison

Chip availability constrains spending; later discussion distinguishes CPU/network supply and a gradual hybrid device/cloud inference trajectory.

FY2025 Q3 | November 19, 2025

QA097 | Supply | Answered

Question Does reduced spending imply lower AI ambition?

Management Management attributed the change to chip availability rather than lower expected token demand.

Follow up Secured supply and delivered capacity.

Questioner: Alex Yao

FY2025 Q4 | Mar 18, 2026

QA098 | CPU | Partially answered

Question How does agentic demand affect cloud pricing amid memory inflation?

Management Management described execution on CPUs, associated memory demand and supply scarcity.

Follow up Specific pricing and cost pass-through.

Questioner: Alex Yao

FY2026 Q1 | May 13, 2026

QA099 | Supply | Answered

Question Are bottlenecks shifting from accelerators to CPUs and networking?

Management Management differentiated China's accelerator constraints and described securing CPU/network supply through established relationships.

Follow up Supplier commitments and continuing accelerator availability.

Questioner: Charlene Liu

FY2026 Q2 | August 12, 2026

QA100 | Inference | Answered

Question What role will on-device inference play?

Management Management expected a gradual hybrid device/cloud approach, shifting more compute to devices as capabilities improve.

Follow up Workload split, device requirements and transition timetable.

Questioner: Alicia a Yap

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Baidu

Cloud | Four quarter comparison

Asset impairments highlight obsolescence risk, while proprietary silicon and utilization support management's subsequent margin argument.

FY2025 Q3 | November 18, 2025

QA085 | Economics | Answered

Question Why impair infrastructure while continuing investment?

Management Management said some existing assets no longer met efficiency needs and described portfolio renewal and utilization improvements.

Follow up Asset-level obsolescence and residual-value assumptions.

Questioner: Wei Xiong

FY2025 Q4 | February 26, 2026

QA086 | Financing | Partially answered

Question How will AI investment be funded and earn returns?

Management Management cited continuing investment, improving financial metrics and leasing or borrowing alternatives.

Follow up Capital cohort returns and investment schedule.

Questioner: Ellie Jiang

FY2026 Q1 | May 18, 2026

QA087 | Inference | Answered

Question How will agents be monetized?

Management Management described application-led value and multiple pricing approaches, expecting value beyond raw tokens.

Follow up Translate application demand into infrastructure requirements.

Questioner: Xiaomeng Zhuang

FY2026 Q2 | August 18, 2026

QA088 | Economics | Answered

Question How can infrastructure margins improve during investment?

Management Management cited GPU-cloud mix, proprietary chips, utilization and operating scale.

Follow up Isolate each contribution and validate full-stack cost.

Questioner: Xiaomeng Zhuang

Transcript links above provide the full exchange. Management claims are not independent performance validation.

CoreWeave

Cloud | Four quarter comparison

Powered-shell execution delays precede continued emphasis on flexible AI infrastructure; CPU/storage requirements for inference remain insufficiently quantified.

FY2025 Q3 | November 10, 2025

QA089 | Capacity | Answered

Question What specifically caused infrastructure delivery delays?

Management Management identified powered-shell delivery rather than power availability for this situation, and described provider diversification.

Follow up Site-specific commissioning milestones; do not generalize its power claim worldwide.

Questioner: Mark Murphy

FY2025 Q4 | February 26, 2026

QA090 | Custom silicon | Partially answered

Question Will inference economics lead to custom ASIC adoption?

Management Management emphasized customer-led NVIDIA demand without committing to a custom-ASIC roadmap.

Follow up Alternative silicon qualification and economics.

Questioner: Mark Murphy

FY2026 Q1 | May 7, 2026

QA091 | Capacity | Partially answered

Question How do component inflation and capacity expansion affect spending?

Management Management emphasized supply-chain execution and spending preceding revenue; the answer did not fully bridge inflation to capacity.

Follow up Price, volume and commissioning bridge.

Questioner: Keith Weiss

FY2026 Q2 | August 11, 2026

QA092 | Inference | Partially answered

Question Can the fleet support inference's additional CPU and storage needs?

Management Management said its infrastructure was built for the full AI cycle and could support inference.

Follow up CPU, storage, network and site retrofit requirements.

Questioner: Raimo Lenschow

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Dell

Systems | Four quarter comparison

Broad cost pressures become a ranked component shortage discussion; later server growth still largely reflects established enterprise customers.

FY2026 Q3 | November 25, 2025

QA101 | Supply | Partially answered

Question How will customers respond to component-driven pricing?

Management Management prioritized supply and product mix amid broad shortages, without quantifying demand elasticity.

Follow up Order elasticity and realized price versus cost.

Questioner: Samik Chatterjee

FY2026 Q4 | February 26, 2026

QA102 | Economics | Partially answered

Question How can AI server economics hold as demand and memory costs rise?

Management Management reaffirmed its operating-margin framework and emphasized broad demand, without isolating memory cost pass-through.

Follow up Quantify component-cost recovery and margin sensitivities.

Questioner: Tim Long

FY2027 Q1 | May 28, 2026

QA103 | Supply | Answered

Question Which constraints matter beyond memory, and in what order?

Management Management identified DRAM, NAND and CPUs first, then hard drives and other components.

Follow up Compare ranked constraints with contracted delivery schedules.

Questioner: Simon Leopold

FY2027 Q2 | September 1, 2026

QA104 | CPU | Answered

Question Is traditional server growth mainly enterprise refresh or AI CPU racks?

Management Management said historical enterprise customers dominate, with emerging neocloud and advanced AI CPU demand.

Follow up Avoid assigning all conventional server growth to agentic AI.

Questioner: Mark Newman

Transcript links above provide the full exchange. Management claims are not independent performance validation.

HPE

Systems | Four quarter comparison

Supply supporting guidance differs from supply meeting demand; price-led server growth must be separated from actual unit delivery.

FY2025 Q4 | December 4, 2025

QA105 | Supply | Partially answered

Question How is the memory cycle reflected in guidance?

Management Management emphasized supply allocation, procurement capability and business mix.

Follow up Memory-price assumptions and sensitivity.

Questioner: Amit Daryanani

FY2026 Q1 | March 9, 2026

QA106 | Supply | Partially answered

Question Is supply sufficient for guidance, and what memory inflation is assumed?

Management Management distinguished supply supporting its outlook from supply insufficient for all customer demand.

Follow up Quantified memory-price assumptions and incremental allocations.

Questioner: Katherine Campagna

FY2026 Q2 | June 1, 2026

QA107 | Memory | Answered

Question Why does storage growth lag servers?

Management Management pointed to migration toward Alletra MP and deferred software revenue within the storage total.

Follow up Normalize platform mix and recognition before comparing growth.

Questioner: Timothy Long

FY2026 Q3 | September 2, 2026

QA108 | Capacity | Answered

Question How much server growth is price versus units, and will supply ease?

Management Management said supply constrained units, with substantial price contribution and stronger units expected while constraints persist.

Follow up Track unit delivery separately from revenue growth.

Questioner: Mark Newman

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Supermicro

Systems | Four quarter comparison

Rack capacity, cooling and component shortages recur; shifting CPU/GPU categories make simple shipment-mix comparisons difficult.

FY2026 Q1 | November 4, 2025

QA109 | Capacity | Partially answered

Question What revenue can manufacturing capacity support?

Management Management gave physical rack capacity and discussed Malaysia, but did not convert it into revenue.

Follow up Rack mix, utilization and accepted production.

Questioner: Ruplu Bhattacharya

FY2026 Q2 | February 3, 2026

QA110 | Cooling | Partially answered

Question Can cooling and system innovations keep arriving ahead of reference designs?

Management Management expressed confidence in earlier integrated delivery based on engineering capability.

Follow up Validated thermals, qualification and actual time-to-operation.

Questioner: Nehal Chokshi

FY2026 Q3 | May 5, 2026

QA111 | Supply | Partially answered

Question How much upside is blocked by CPU and other shortages?

Management Management confirmed CPU, memory, storage and some GPU constraints without sizing foregone revenue.

Follow up Component-specific unmet demand and timing.

Questioner: Asiya Merchant

FY2026 Q4 | August 11, 2026

QA112 | CPU | Partially answered

Question How does inference change CPU/GPU shipment mix?

Management Management offered an approximate mix while noting blurred categories for agentic and edge GPUs.

Follow up Consistent unit/revenue definitions and forward mix.

Questioner: W. Chiu

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Lenovo

Systems | Four quarter comparison

Enterprise inference and cooling guide portfolio changes; management explicitly clarifies that pipeline is not committed orders.

FY2026 Q2 | November 20, 2025

QA159 | Cooling | Answered

Question How does liquid cooling support the infrastructure opportunity?

Management Management linked Neptune to energy efficiency and air-to-liquid conversion services.

Follow up Customer-validated energy and deployment results.

Questioner: Jim Au, DBS; read by Lixi Yuan

FY2026 Q3 | February 12, 2026

QA160 | Inference | Answered

Question How does restructuring position Lenovo for inference?

Management Management described portfolio simplification, stronger inference offerings and changes to enterprise delivery and sales.

Follow up Verify production adoption and operational savings.

Questioner: Leping, Huatai Securities; read by Lixi Yuan

FY2026 Q4 | May 22, 2026

QA161 | Inference | Partially answered

Question What customer and training/inference mix sits in the pipeline?

Management Management described broad geographic/customer participation and expected a shift toward inference without a numerical split.

Follow up Pipeline mix and conversion by workload.

Questioner: William Huang, Huatai Securities; read by Lixi Yuan

FY2027 Q1 | August 13, 2026

QA162 | Capacity | Partially answered

Question Is the AI pipeline committed orders, and when will it convert?

Management Management explicitly distinguished uncommitted or unqualified opportunities from orders and did not give a conversion forecast.

Follow up Do not add pipeline to backlog; obtain conversion and delivery schedules.

Questioner: Albert Hung, JPMorgan; read by Lixi Yuan

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Arista

Network | Four quarter comparison

AI network boundaries blur, but back-end urgency, shipment timing and alternative-component qualification shape near-term delivery.

FY2025 Q3 | November 5, 2025

QA113 | Network | Answered

Question How does AI back-end traffic affect front-end upgrades?

Management Management said accelerator networks increasingly pressure general cloud networking, making revenue boundaries harder to separate.

Follow up Measure traffic and upgrade needs rather than treating front/back-end categories as fixed.

Questioner: Meta Marshall

FY2025 Q4 | February 12, 2026

QA114 | Capacity | Answered

Question Why does network growth not track hyperscaler capex immediately?

Management Management cited deployment sequence, product qualification and shipment timing across years.

Follow up Network acceptance and delivery schedule.

Questioner: Samik Chatterjee

FY2026 Q1 | May 5, 2026

QA115 | Network | Partially answered

Question How does agentic AI change front-end versus back-end demand?

Management Management described strong immediate back-end demand and developing distributed inference and enterprise use.

Follow up A stable traffic ratio and timing of front-end refresh.

Questioner: Sean O'Loughlin

FY2026 Q2 | August 4, 2026

QA116 | Supply | Answered

Question How are upstream decommitments and shortages being addressed?

Management Management described supplier engagement and qualifying alternatives in memory, boards and optics while expecting persistent industry constraints.

Follow up Validate alternate-component qualification and firm commits.

Questioner: Michael Ng

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Cisco

Network | Four quarter comparison

Edge use cases and custom networking silicon combine with memory redesign and direct supplier engagement; demand remains lumpy.

FY2026 Q1 | November 12, 2025

QA117 | Inference | Answered

Question Where and how will enterprise edge inference deploy?

Management Management emphasized latency-sensitive local use cases and multiple deployment models.

Follow up Workload-specific latency budget and economics versus centralized execution.

Questioner: Antoine Chkaiban

FY2026 Q2 | February 11, 2026

QA118 | Capacity | Answered

Question Why do AI orders not rise directly with customer capex?

Management Management pointed to lumpy orders concentrated in a small customer set and incremental use-case wins.

Follow up Compare firm orders and delivery dates instead of capex proxies.

Questioner: David Vogt

FY2026 Q3 | May 13, 2026

QA119 | Supply | Answered

Question How is Cisco securing components beyond memory?

Management Management cited silicon commitments, memory-reduction redesigns, migration to newer memory and supplier agreements.

Follow up Distinguish portfolio redesign benefits from AI networking specifics.

Questioner: Aaron Rakers

FY2026 Q4 | August 12, 2026

QA120 | Custom silicon | Answered

Question How does Silicon One change control of supply and lead times?

Management Management reported no significant lead-time issues and cited direct TSMC coordination and reduced dependence on merchant-silicon allocations.

Follow up Compare actual deliveries and supplier commitments across product families.

Questioner: Ben Bohlen

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Vertiv

Power cooling | Four quarter comparison

Thermal and electrical integration grows more important, and larger projects introduce internal logistics and execution dependencies.

FY2025 Q3 | October 22, 2025

QA125 | Capacity | Partially answered

Question Can capacity expansion keep up with backlog?

Management Management emphasized existing-site expansion and productivity, declining specific next-year capex.

Follow up Capacity commissioning and throughput by product.

Questioner: Andrew Alec Kaplowitz

FY2025 Q4 | February 11, 2026

QA126 | Cooling | Answered

Question How will thermal architectures change mix and content?

Management Management expected hybrid thermal chains and continued CDU relevance, dependent on workload and resiliency.

Follow up Validate full-site thermal design rather than extrapolating one GPU configuration.

Questioner: Mark Trevor Delaney

FY2026 Q1 | April 22, 2026

QA127 | Platform | Answered

Question Is procurement moving from individual products to integrated systems?

Management Management described a gradual, partial shift toward modules and converged systems with varying customer adoption.

Follow up Define integration responsibility and interface ownership.

Questioner: Noah Kaye

FY2026 Q2 | July 29, 2026

QA128 | Capacity | Partially answered

Question What caused delivery timing shifts and congestion?

Management Management cited project complexity and internal as well as external supply interdependencies.

Follow up Specific delayed milestones and recurrence risk.

Questioner: Scott Davis

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Eaton

Power cooling | Four quarter comparison

Higher-voltage DC and cooling expand the discussion from components to complete capabilities, including protection and service.

FY2025 Q3 | November 4, 2025

QA121 | Power | Partially answered

Question How does higher-voltage DC change content per megawatt?

Management Management explained fewer conversions and its DC capabilities, but the future architecture's content remained conditional.

Follow up Architecture-specific bill of materials and conversion losses.

Questioner: Andrew Kaplowitz

FY2025 Q4 | February 3, 2026

QA122 | Cooling | Partially answered

Question How do cooling changes affect Eaton's opportunity?

Management Management emphasized growth and power/cooling integration and cited evolving accessible content.

Follow up Separate market-content estimates from actual deployed revenue.

Questioner: Andrew Obin

FY2026 Q1 | May 5, 2026

QA123 | Cooling | Partially answered

Question Can Boyd integrate and expand capacity without disrupting delivery?

Management Management described deliberate integration and historically modest capital intensity.

Follow up Specific capacity additions, dates and execution risks.

Questioner: Joseph Ritchie

FY2026 Q2 | July 31, 2026

QA124 | Power | Answered

Question How does the DC transition affect Eaton's content and position?

Management Management confirmed its current content estimate and identified transformers, DC breakers, power electronics, cooling and service as required capabilities.

Follow up Future architecture content and efficiency still require design-level validation.

Questioner: Andrew Kaplowitz

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Schneider Electric

Power cooling | Four quarter comparison

Generally normalized equipment lead times coexist with a gradual DC architecture transition; broad electrification also drives growth.

FY2025 Q3 | October 30, 2025

QA163 | Capacity | Partially answered

Question Is data-center equipment capacity limiting growth?

Management Management said lead times had generally normalized and capacity investment would continue.

Follow up Utilization was not quantified; portfolio conditions differ.

Questioner: Phil Buller, JPMorgan

FY2025 Q4 | February 26, 2026

QA149 | Power | Partially answered

Question How is Schneider positioned for higher-voltage DC, and what is the data-center backlog?

Management Management discussed evolving infrastructure requirements but did not provide the requested backlog breakdown.

Follow up Qualified system readiness and comparable backlog disclosure.

Questioner: Andre Kukhnin

FY2026 Q1 | April 30, 2026

QA164 | Power | Answered

Question What is the DC-power architecture roadmap?

Management Management described a gradual rollout, early sidecar solutions followed by more integrated systems and selected partners.

Follow up Qualification dates and integration responsibility.

Questioner: Gael de-Bray, Deutsche Bank

FY2026 Q2 | July 30, 2026

QA150 | Power | Answered

Question Is strong demand mainly data centers and semiconductors?

Management Management described broader electrification, grid and industrial demand alongside data centers.

Follow up Avoid assigning all electrical growth to AI.

Questioner: Benedict Uglow

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Equinix

Colocation | Four quarter comparison

Inference interconnection and rack density increase the importance of site-specific power; contracted and expected supply must be separated.

FY2025 Q3 | October 29, 2025

QA131 | Power | Partially answered

Question How secure is power for new projects?

Management Management described contracted or advanced-stage power arrangements and utility-related investment requirements.

Follow up Separate executed agreements from ongoing discussions.

Questioner: James Schneider

FY2025 Q4 | February 11, 2026

QA132 | Inference | Answered

Question When is enterprise inference arriving and why does interconnection matter?

Management Management cited emerging enterprise demand and private connectivity as a deployment enabler.

Follow up Workload-level demand attribution and connection economics.

Questioner: Cameron McVeigh

FY2026 Q1 | Apr 29, 2026

QA133 | Power | Answered

Question Does high-density inference face space, thermal or power limits?

Management Management identified power as the largest constraint and said densification can require reserving adjacent space.

Follow up Usable power and service-level compliance by facility.

Questioner: Joseph Osha

FY2026 Q2 | July 29, 2026

QA134 | Power | Partially answered

Question How much expansion has an executed utility agreement?

Management Management grouped contracted power with high-confidence future contracting instead of giving the requested proportion.

Follow up Executed-versus-expected power split and energization gates.

Questioner: Michael Elias

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Digital Realty

Colocation | Four quarter comparison

Inference demand spans a wide range of deployment sizes; development runway is not a dated schedule of energized capacity.

FY2025 Q3 | October 23, 2025

QA165 | Inference | Partially answered

Question Is latency-sensitive AI demand appearing in major metros?

Management Management described broad AI/cloud dialogue and forthcoming capacity, without specific latency-sensitive application examples.

Follow up Named workload requirements and actual latency budgets.

Questioner: Jonathan Petersen

FY2025 Q4 | February 5, 2026

QA166 | Inference | Answered

Question Are larger enterprise blocks a sign of inference demand?

Management Management described emerging conversations for larger contiguous blocks as inference develops.

Follow up Distinguish discussions from signed, operating inference deployments.

Questioner: Michael Elias

FY2026 Q1 | April 23, 2026

QA129 | Inference | Answered

Question What do different lease sizes reveal about AI demand?

Management Management cited inference-related demand across very large and sub-megawatt deployments.

Follow up Do not map all inference to small edge sites.

Questioner: Michael Rollins

FY2026 Q2 | July 23, 2026

QA130 | Power | Partially answered

Question What power is committed and when does development capacity arrive?

Management Management discussed development and preleasing without a complete year-by-year energized-power schedule.

Follow up Utility commitments and energization dates by site.

Questioner: Michael Rollins

Transcript links above provide the full exchange. Management claims are not independent performance validation.

ASML

Equipment | Four quarter comparison

AI announcements cannot be mechanically converted into tool demand; volume, process intensity, reuse and productivity all matter.

FY2025 Q3 | Oct 15, 2025

QA138 | Foundry | Partially answered

Question Can AI capacity announcements be translated into EUV demand?

Management Management declined a mechanical conversion and emphasized broader customer participation.

Follow up Firm wafer plans, layer counts and tool productivity.

Questioner: Didier Scemama

FY2025 Q4 | Jan 28, 2026

QA139 | Foundry | Partially answered

Question Will DRAM architecture changes reduce EUV demand through tool reuse?

Management Management described multi-node planning and process simplification benefits from EUV.

Follow up Net tool additions after reuse and productivity.

Questioner: Didier Scemama

FY2026 Q1 | Apr 15, 2026

QA140 | Capacity | Answered

Question Will EUV equipment become the expansion bottleneck?

Management Management pointed to manufacturing, cycle-time and productivity levers and said it was not currently the bottleneck.

Follow up Compare promised throughput with customer delivery windows.

Questioner: Lee Simpson

FY2026 Q2 | July 15, 2026

QA141 | Memory | Partially answered

Question How much memory equipment growth is volume versus lithography intensity?

Management Management identified both HBM/DDR volumes and additional layers without quantifying the split.

Follow up Volume/intensity bridge; equipment revenue does not equal bit-capacity growth.

Questioner: Stéphane Houri

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Applied Materials

Equipment | Four quarter comparison

The discussion links packaging and device transitions to existing-fab productivity; segment and capacity-growth decompositions remain incomplete.

FY2025 Q4 | November 13, 2025

QA135 | Capacity | Answered

Question Are memory customers constrained by fab-shell space?

Management Management said aggregate space appeared available but customer-specific installation and qualification schedules differed.

Follow up Check each customer rather than treating macro capacity as immediately usable.

Questioner: Atif Malik

FY2026 Q1 | Feb 12, 2026

QA167 | Packaging | Partially answered

Question How fast will advanced packaging grow relative to equipment overall?

Management Management highlighted HBM and 3D stacking plus larger future substrates without quantifying relative growth.

Follow up Packaging growth decomposition and customer qualification.

Questioner: Krish Sankar, TD Cowen

FY2026 Q2 | May 14, 2026

QA136 | Foundry | Answered

Question Is improving existing-fab output and yield an incremental driver?

Management Management confirmed customer emphasis on productivity and yield while new space takes time.

Follow up Quantify good-wafer output improvements.

Questioner: Harlan Sur

FY2026 Q3 | August 13, 2026

QA137 | Foundry | Partially answered

Question Which segment leads incremental equipment strength?

Management Management described broad AI-related demand across leading logic, DRAM and packaging without ranking them.

Follow up Segment-specific incremental capacity and growth.

Questioner: Jim Schneider

Transcript links above provide the full exchange. Management claims are not independent performance validation.

Lam Research

Equipment | Four quarter comparison

Facility timing, NAND upgrades and broader packaging adoption shape demand; planning visibility is not equivalent to binding orders.

FY2026 Q1 | October 22, 2025

QA146 | Capacity | Answered

Question What limits acceleration of semiconductor supply?

Management Management said physical facility timing can exceed equipment lead times.

Follow up Customer-specific cleanroom availability and tool acceptance.

Questioner: Harlan Sur

FY2026 Q2 | January 28, 2026

QA147 | Memory | Answered

Question Are NAND migrations slowing, and when does new capacity arrive?

Management Management distinguished DRAM prioritization from continuing NAND upgrades, with capacity additions following.

Follow up Separate technology upgrades from net wafer capacity.

Questioner: Atif Malik

FY2026 Q3 | April 22, 2026

QA148 | Supply | Partially answered

Question Do customer memory agreements create equipment commitments?

Management Management described longer planning visibility and operational benefits without confirming equivalent down-payment obligations.

Follow up Binding equipment orders versus planning discussions.

Questioner: Vivek Arya

FY2026 Q4 | Jul 29, 2026

QA168 | Packaging | Answered

Question What drives the stronger advanced-packaging outlook?

Management Management cited broad HBM and logic packaging adoption, with panel packaging still early.

Follow up Separate technology adoption, pull-ins and sustained volume.

Questioner: Harlan Sur, JPMorgan

Transcript links above provide the full exchange. Management claims are not independent performance validation.

KLA

Equipment | Four quarter comparison

Logic, DRAM and packaging create process-control demand; pricing, technology intensity and physical capacity require separate treatment.

FY2026 Q1 | October 29, 2025

QA142 | Packaging | Partially answered

Question Does packaging growth include HBM as well as logic?

Management Management expressed confidence in both markets without separating contributions.

Follow up HBM-specific inspection demand and share.

Questioner: Brian Chin

FY2026 Q2 | January 29, 2026

QA143 | Memory | Answered

Question How does equipment growth differ between DRAM and logic?

Management Management expected faster DRAM growth with both HBM and conventional-memory demand.

Follow up Track actual equipment deliveries and qualified output.

Questioner: Chris Caso

FY2026 Q3 | April 29, 2026

QA144 | Capacity | Partially answered

Question Is incremental demand from new fabs, migration or yield work?

Management Management described urgent near-term deliveries and broad future greenfield activity, without a precise decomposition.

Follow up Separate capacity additions from technology transitions.

Questioner: Harlan Sur

FY2026 Q4 | July 28, 2026

QA145 | Economics | Partially answered

Question Are tool price changes cost recovery or value pricing?

Management Management described both costs and delivered value, with new-product and existing-product mechanisms differing.

Follow up Like-for-like pricing, timing and margin bridge.

Questioner: Sreekrishnan Sankarnarayanan

Transcript links above provide the full exchange. Management claims are not independent performance validation.