



AI Infrastructure Architecture Evolution

2015-2026

Silicon generations · Memory and interconnects
Conference disclosures · Systems and software

A generation-by-generation synthesis of conference disclosures and primary technical sources.
Evidence cutoff: 28 September 2026.
Complete English edition, paired with a full Simplified Chinese report.

Contents

01 Scope and executive synthesis	3
02 Bandwidth, status and topology conventions	3
03 Tomahawk1-4: radix and SerDes drive the system	3
04 Tomahawk5, Tomahawk6 and Tomahawk Ultra	4
05 Jericho: scheduled fabrics and regional scale	4
06 NICs, SmartNICs and complementary connectivity	5
07 Generation and network comparison tables	6
08 Custom XPU's and portfolio boundaries	9
09 Integrated stack and software	12
10 Derived metrics and architectural trade-offs	17
11 Conference index and naming crosswalk	18
Source map	20
References	21

01 Scope and executive synthesis

Deep coverage: Tomahawk, Jericho, AI Ethernet NICs and their system roles. Full coverage: SerDes, CPO and publicly named custom-XPU programs. CPU and accelerator schemas are retained, but Broadcom is not treated as a merchant server-CPU or general-purpose GPU vendor. Period: 2015-28 September 2026. The latest reviewed events include ISSCC 2025, Hot Chips 2025, the public Hot Chips 2026 program and 2026 partner announcements. Storage connectivity is proportional; VMware is outside the silicon architecture scope.

Broadcom's evolution is from supplying switch capacity to shaping the entire communication path of an AI system. High radix reduces network tiers; NIC transport handles path diversity and recovery; scheduled fabrics control contention; optics changes physical reach and power. The custom-XPU business connects these technologies to customer-specific compute. The important architectural comparison is not simply Ethernet versus a proprietary fabric, but which layer owns scheduling, ordering, congestion and failure recovery.

02 Bandwidth, status and topology conventions

Switch capacity is quoted in the vendor's native convention and is not automatically the sum of both directions. Endpoint Ethernet rates are line rates; dividing Gb/s by eight gives nominal GB/s before protocol overhead. Jericho front-panel bandwidth and fabric-interface bandwidth must be separated: adding them can produce a headline that is not usable endpoint injection bandwidth. Chip availability, OEM switch qualification and deployed system scale are separate milestones. n/d means not established in reviewed sources; n/a means not applicable.

03 Tomahawk1-4: radix and SerDes drive the system

Tomahawk established the high-bandwidth cloud-switch line in the 2015-era baseline. Tomahawk2 doubled capacity to 6.4 Tb/s; Tomahawk3 doubled it again to 12.8 Tb/s and enabled dense 400GbE configurations through 50G PAM4 signaling. Tomahawk4 reached 25.6 Tb/s and subsequently entered production. These transitions are not just larger crossbars: each changes the number of endpoints that can be connected at a fixed radix, the number of tiers, cable count and optical power.

Architectural interpretation: at fixed total switching bandwidth, increasing endpoint speed reduces the number of endpoint-facing ports. Capacity doubling can therefore be consumed entirely by speed doubling, leaving radix unchanged. A system designer should ask whether a new generation buys more nodes per tier, higher bandwidth per node or lower oversubscription. Those are different benefits. Port counts must also reserve uplinks; a 64-port switch is not automatically a nonblocking 64-endpoint leaf.

04 Tomahawk5, Tomahawk6 and Tomahawk Ultra

Tomahawk5 couples 51.2 Tb/s switching with 100G-class PAM4 lanes. The ISSCC 2025 paper 16.1 identifies a monolithic 5 nm implementation with 512 PAM4 SerDes lanes; the reviewed press kit establishes that implementation mapping, while product documents establish supported port configurations. This is a useful example of conference synthesis: the circuit venue explains realization and power-management concerns, while the product source explains the network building block. The press kit is not treated as the full paper.

Tomahawk6 reaches 102.4 Tb/s and offers 100G- and 200G-class SerDes configurations. Its architecture combines higher radix with routing, congestion and telemetry functions for AI. The 2025 announcement and the 2026 production announcement are recorded separately. CPO variants are not automatically identical to the electrical package: optical-engine integration changes the package, cooling, serviceability and link-budget problem even when the logical switch capacity remains the same.

Tomahawk Ultra is a specialization branch, not simply the next numbered Tomahawk. The Hot Chips 2025 presentation emphasizes 51.2 Tb/s operation for small packets, roughly 250 ns switching latency, link-level retry, shortened headers and in-network collectives. These features address scale-up traffic where transaction overhead and latency can dominate. A small-packet throughput claim and a bulk-bandwidth claim are different tests; full-size frame line rate cannot establish the former.

05 Jericho: scheduled fabrics and regional scale

The Jericho lineage targets richer traffic management and modular fabric construction. Jericho2 combines packet processing, buffering and fabric attachment; its 9.6 Tb/s headline includes interfaces that must not all be counted as endpoint traffic. Jericho3-AI extends this approach to AI fabrics using coordinated scheduling and a separate fabric element. Its significance is contention management across a distributed switch system, rather than a claim that every individual chip has the capacity of the whole cluster.

Jericho4 adds 3 nm implementation, 200G PAM4, MACsec and the 3.2 Tb/s HyperPort abstraction formed from four 800GbE links. The launch discusses RoCE over regional distances beyond 100 km. The availability paragraph says sampling, despite a headline using “ships”; this report uses sampling unless a later production source establishes otherwise. Longer reach expands placement freedom but cannot remove propagation latency, so the suitable parallelism and synchronization cadence remain workload-dependent.

06 NICs, SmartNICs and complementary connectivity

Stingray represents infrastructure offload: programmable processing moves cloud services away from the host CPU. Thor and Thor Ultra represent the endpoint communication path. Thor Ultra's PCIe 6 ×16 interface, 800G Ethernet, multipath delivery, out-of-order placement, selective retransmission and programmable congestion control target AI collective traffic. PSP offload and device security complement transport. These functions do not make it interchangeable with a DPU running arbitrary infrastructure services.

SerDes, retimers, optical DSPs, CPO, PCIe switches and storage connectivity are adjacent but distinct product layers. A retimer restores an electrical link; an optical DSP conditions a modulated optical channel; a switch routes traffic. Their power and latency belong in the end-to-end path rather than being hidden in a switch-chip number. OCP and OFC disclosures are especially relevant for these integration boundaries, while OEM demonstrations establish system compatibility rather than universal deployment.

07 Generation and network comparison tables

Switch generations and capacity basis

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Family	First disclosure / production	Capacity	Implementation / role	Status
Tomahawk1	2014-2015	3.2 Tb/s	Cloud Ethernet	Historical shipping
Tomahawk2	2016	6.4 Tb/s	64 × 100GbE	Historical shipping
Tomahawk3	2017 / 2018	12.8 Tb/s	50G PAM4	Historical shipping
Tomahawk4	2019 / 2020	25.6 Tb/s	High-radix Ethernet	Shipping / available
Tomahawk5	2022 / 2023	51.2 Tb/s	5 nm; 512 PAM4 lanes	Shipping / available
Tomahawk6	2025 / 2026	102.4 Tb/s	100G / 200G PAM4	Shipping / available
Tomahawk Ultra	2025	51.2 Tb/s	Small-packet / low latency	Shipping / available
Jericho2	2018 / 2019	4.8 Tb/s packet path; separate fabric I/O	Buffered modular switch	Historical shipping
Jericho3-AI	2023	28.8 Tb/s vendor combined basis	Scheduled AI fabric	Announced; GA n/d
Jericho4	2025	3.2 Tb/s per HyperPort; chip total n/d	3 nm; 200G PAM4	Sampling; GA n/d

Networking fixed schema · 1/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	GA / milestone	Status	Class
Tomahawk5	2023	Shipping / available	Scale-out switch
Tomahawk Ultra	2025	Shipping / available	Scale-up switch
Tomahawk6	2026 production	Shipping / available	Scale-up / out switch
Jericho4	2025 sampling	Sampling; GA n/d	Fabric router
Thor Ultra	2025 announced	Announced; GA n/d	AI NIC
Stingray	2020 deployment	Historical shipping	SmartNIC

Networking fixed schema · 2/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Ports × speed	SerDes	Host interface	Programmable engine
Tomahawk5	64 × 800GbE	100G PAM4	n/a	Switch pipeline
Tomahawk Ultra	64 × 800GbE	100G PAM4	n/a	Switch / collectives
Tomahawk6	102.4 Tb/s total	100G / 200G PAM4	n/a	Switch pipeline
Jericho4	HyperPort 4 × 800G	200G PAM4	n/a	Traffic manager
Thor Ultra	800 Gb/s	100G / 200G PAM4	PCIe 6 ×16	Programmable congestion pipeline
Stingray	n/d	n/d	PCIe; n/d	Programmable offload

Networking fixed schema · 3/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Embedded CPU	On-card memory	Transport / RDMA	Congestion control
Tomahawk5	n/d	n/d	Ethernet	Cognitive Routing
Tomahawk Ultra	n/d	n/d	Ethernet / LLR	Adaptive routing
Tomahawk6	n/d	n/d	Ethernet	Cognitive Routing 2.0
Jericho4	n/d	n/d	Ethernet / RoCE	Deep-buffer fabric
Thor Ultra	n/d	n/d	RoCE / UEC	Sender / receiver programmable
Stingray	n/d	n/d	n/d	n/d

Networking fixed schema · 4/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Key offloads	Power
Tomahawk5	L2/L3 / telemetry	n/d
Tomahawk Ultra	Header reduction / collectives	n/d
Tomahawk6	Telemetry / CPO option	n/d
Jericho4	MACsec	n/d
Thor Ultra	PSP / OOO / retry	n/d
Stingray	Cloud infrastructure	n/d

08 Custom XPU's and portfolio boundaries

Publicly named collaborations establish Broadcom's role in customer-specific accelerators. Meta identifies the MTIA collaboration, and the June 2026 OpenAI announcement identifies Jalapeño engineering samples and Broadcom's implementation and networking contribution. Customer architecture ownership and supplier implementation responsibility must remain distinct. There is no single public Broadcom XPU instruction set or uniform accelerator specification that can be assigned to every customer.

Architectural interpretation: reusable SerDes, memory interfaces, die-to-die links and physical-design capabilities can shorten implementation without making customer compute architectures identical. The software burden remains attached to the customer's compiler, runtime and model kernels. A useful cross-vendor comparison should place MTIA specifications in the Meta report and only documented supplier contributions here, preventing double counting of both engineering claims and installed capacity.

Host CPU boundary · 1/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Codename	GA	Status	Core µarch	Cores / threads
----------------------------	----------	----	--------	------------	-----------------

No in-scope merchant host CPU	n/a	n/a	n/a	n/a	n/a
-------------------------------	-----	-----	-----	-----	-----

Host CPU boundary · 2/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Process per die	Dies / packaging	L2 per core	L3 / domain
----------------------------	-----------------	------------------	-------------	-------------

No in-scope merchant host CPU	n/a	n/a	n/a	n/a
-------------------------------	-----	-----	-----	-----

Host CPU boundary · 3/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Memory: channels; rate; GB/s	Sockets / links	PCIe / CXL	TDP	Vector / matrix ISA
----------------------------	------------------------------	-----------------	------------	-----	---------------------

No in-scope merchant host CPU	n/a	n/a	n/a	n/a	n/a
-------------------------------	-----	-----	-----	-----	-----

Customer-owned accelerator programs · 1/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Architecture	GA / milestone	Status	Process	Dies / packaging
---------	--------------	----------------	--------	---------	------------------

MTIA collaboration	Customer-specific	n/d	See Meta report	n/d	n/d
Jalapeño	OpenAI inference	n/d	Sampling; GA n/d	n/d	n/d

Customer-owned accelerator programs · 2/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Compute units	Dense BF16	Dense FP8	Dense FP6 / FP4	FP64 vector / matrix
MTIA collaboration	n/d	n/d	n/d	n/d	n/d
Jalapeño	n/d	n/d	n/d	n/d	n/d

Customer-owned accelerator programs · 3/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Memory: type; stacks; GB; TB/s	LLC / SRAM
MTIA collaboration	n/d	n/d
Jalapeño	n/d	n/d

Customer-owned accelerator programs · 4/4


n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Scale-up: links; GB/s per direction; domain	Host link	Power / cooling	Form factor
MTIA collaboration	n/d	n/d	n/d	n/d
Jalapeño	n/d	n/d	n/d	n/d

09 Integrated stack and software


The deployed network combines silicon SDKs, operating systems such as SONiC in supported platforms, management and telemetry, and endpoint transport software. Open protocols do not guarantee identical behavior across vendors: feature negotiation, congestion tuning and firmware maturity can dominate interoperability. UEC compliance, RoCE compatibility and scale-up memory semantics are separate questions. Network software must be qualified with the target accelerator collective library and failure scenarios.

2015-2020 network stack



Product roles, not a claim that all families are required in one deployment.

2023-2026 network stack



Product roles, not a claim that all families are required in one deployment.

Interconnect roles · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Interconnect	Year / generation	Lane rate	Lanes / link	GB/s per link per direction
Ethernet / SUE	2025-2026	100G / 200G class	n/d	n/d
HyperPort	2025	n/d	4 × 800GbE	400 GB/s nominal

Interconnect roles · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Interconnect	Links / device	Topology	Domain limit	Coherence / semantics
Ethernet / SUE	n/d	Switched	Topology-dependent	Transport does not establish CPU coherence
HyperPort	n/d	Logical grouped port	n/d	Ethernet / RoCE

System integration boundary · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Platform	Year / milestone	CPU / accelerator count	Scale-up domain
OEM switch platforms	2015-2026	External CPUs / XPU	Topology-dependent

System integration boundary · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Platform	Scale-out per accelerator	Rack power	Cooling
OEM switch platforms	n/d	n/d	Platform-dependent

Software integration contracts

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Layer	Interface	Function	Boundary
Switch control	SDK / supported open interfaces	Forwarding and telemetry	ASIC-specific implementation
Infrastructure	Firmware / driver / acceleration APIs	Packet, storage and security offload	SKU and platform dependent
System integration	Host software and partner stack	Resource and traffic management	No single vendor-only AI framework

10 Derived metrics and architectural trade-offs

Nominal bandwidth normalization

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Quantity	Calculation	Meaning
800GbE	$800/8 = 100 \text{ GB/s}$	Per direction before overhead
3.2T HyperPort	$4 \times 800/8 = 400 \text{ GB/s}$	Logical port; not chip total
Tomahawk1 → 6	$102.4/3.2 = 32 \times$	Native switch-capacity convention
HBM / FLOP; performance / W	n/a; n/d	No uniform compute SKU / power basis

The main trade-off is where complexity sits. An endpoint-scheduled fabric places congestion response in NICs and hosts; a fabric-scheduled design moves more coordination into the network. Low-latency scale-up adds requirements for small packets, retry and collective support that differ from conventional cloud forwarding. CPO shortens electrical reach and can reduce I/O energy, but places optical reliability and maintenance closer to expensive switch silicon. None of these choices eliminates the need to measure job completion time under contention and failure.

11 Conference index and naming crosswalk

Verified disclosure map · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Year	Paper / talk / disclosure
ISSCC	2025	16.1 Tomahawk5: 51.2Tb/s 5nm Monolithic Switch Chip for AI/ML Networking
Hot Chips	2025	Tomahawk Ultra
Hot Chips	2026	Thor Ultra: An Ethernet NIC Chip Optimized for AI & HPC
OCP Global Summit	2025	End-to-end AI networking

Verified disclosure map · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Presenter / authors	Product	Evidence type	Source
ISSCC	Broadcom	Tomahawk5	Implementation; press kit reviewed	C01
Hot Chips	Broadcom	Tomahawk Ultra	Architecture slides	C02
Hot Chips	Hemal Shah	Thor Ultra	Program verified; vendor specs used	C03 E03
OCP Global Summit	Broadcom	Switch / NIC / CPO	System demonstrations	E04

ISCA, MICRO, HPCA and ASPLOS are not used here to manufacture a nonexistent product-paper lineage. For this portfolio, the strongest identified bridge is Hot Chips architecture plus ISSCC implementation and vendor/OCP system documentation. OFC is a natural follow-up channel for optics; the present bibliography does not claim exhaustive coverage of every OFC or DAC session. Product names and codenames are crosswalked only where the vendor identifies them.

Product names

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Marketing family	Product identifier	Distinction
Tomahawk1	BCM56960	Cloud switch baseline
Tomahawk5	BCM78900	General high-radix switch
Tomahawk Ultra	BCM78920	Scale-up specialization
Jericho3-AI	BCM88890	Scheduled AI fabric
Thor Ultra	n/d	AI NIC; not Tomahawk Ultra

The remaining gaps are quantitative power at a fixed port configuration, full buffer organization for every generation, per-lane electrical detail and customer-XPU implementation specifications. These constrain direct energy comparisons. The most useful next tests would hold endpoint count, port speed, oversubscription, collective mix and failure load constant, then measure delivered bandwidth and tail job latency across endpoint- and fabric-scheduled alternatives.

Source map

- 01 Scope and executive synthesis - C01, C02, C03, E01, E02, E03, E11, E12
- 02 Bandwidth, status and topology conventions - E01, E03, P04, P05
- 03 Tomahawk1-4: radix and SerDes drive the system - E01, E05, E06, E07, P01, P03
- 04 Tomahawk5, Tomahawk6 and Tomahawk Ultra - C01, C02, E01, E04, E08, E09, P01, P02
- 05 Jericho: scheduled fabrics and regional scale - E02, E13, P04, P05
- 06 NICs, SmartNICs and complementary connectivity - E03, E04, E10
- 07 Generation and network comparison tables - C01, C02, E01, E02, E03, E05, E06, E07, E08, E09, E10, E13, P01, P03, P04, P05
- 08 Custom XPU's and portfolio boundaries - E11, E12
- 09 Integrated stack and software - C02, E01, E02, E03, E04, E10, P01, P02
- 10 Derived metrics and architectural trade-offs - C02, E01, E02, E03, E04, E11, E12, E13, P03
- 11 Conference index and naming crosswalk - C01, C02, C03, E03, E04, E13, P01, P02, P03, P05

References

Product and technical documentation

[P01] Tomahawk5 BCM78900

<https://www.broadcom.com/products/ethernet-connectivity/switching/strataxgs/bcm78900-series>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P02] Tomahawk Ultra BCM78920

<https://www.broadcom.com/products/ethernet-connectivity/switching/strataxgs/bcm78920-series>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P03] Tomahawk1 BCM56960

<https://www.broadcom.com/products/ethernet-connectivity/switching/strataxgs/bcm56960-series>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P04] Jericho2 datasheet

<https://docs.broadcom.com/doc/BCM88690-9-6-Tb/s-Integrated-Packet-Processor-Traffic-Manager-and-Fabric-Interface-Single-Chip-Device-DS>

Primary source; accessed by 2026-09-28

[P05] Jericho3-AI product

<https://www.broadcom.com/products/ethernet-connectivity/switching/stratadnx/bcm88890>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

Conference records and implementation disclosures

[C01] ISSCC 2025 press kit; paper 16.1 Tomahawk5

<https://submissions.mirasmart.com/ISSCC2025/PDF/ISSCC2025PressKit.pdf>

Official conference program or press material; not the full paper.

[C02] Tomahawk Ultra, Hot Chips 2025

<https://hc2025.hotchips.org/assets/program/conference/day1/TU-HotChips-2025-Final-2025-08-25-v1.pdf>

Primary source; accessed by 2026-09-28

[C03] Thor Ultra, Hot Chips 2026 program

<https://hc2026.hotchips.org/program/conference/>

Official indexed record or related announcement; full document download unavailable.

Launches, ecosystem and deployment

[E01] Tomahawk6 launch

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-ships-tomahawk-6-worlds-first-1024-tbps-switch>

Primary source; accessed by 2026-09-28

[E02] Jericho4 launch

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-ships-jericho4-enabling-distributed-ai-computing-across>

Primary source; accessed by 2026-09-28

[E03] Thor Ultra 800G AI Ethernet NIC

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-introduces-industrys-first-800g-ai-ethernet-nic>

Primary source; accessed by 2026-09-28

[E04] OCP 2025 end-to-end AI networking

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-delivers-future-ai-infrastructure-end-end-ai-networking>

Primary source; accessed by 2026-09-28

[E05] Tomahawk2 announcement

<https://www.broadcom.com/news/product-releases/broadcom-first-to-deliver-64-ports-of-100ge-with-tomahawk-ii-ethernet-switch>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E06] Tomahawk3 production

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-achieves-mass-production-industry-leading-128-tbps>

Primary source; accessed by 2026-09-28

[E07] Tomahawk4 production

<https://www.broadcom.com/company/news/product-releases/53966>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E08] Tomahawk5 production

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-now-shipping-worlds-first-512-tbps-switch-production>

Primary source; accessed by 2026-09-28

[E09] Tomahawk6 production

<https://www.broadcom.com/company/news/product-releases/64031>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E10] Stingray SmartNIC deployment

<https://www.broadcom.com/company/news/product-releases/53106>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E11] Meta custom silicon partnership

<https://www.broadcom.com/company/news/product-releases/64236>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E12] OpenAI Jalapeno engineering samples

<https://investors.broadcom.com/news-releases/news-release-details/openai-and-broadcom-unveil-llm-optimized-intelligence-processor>

Primary source; accessed by 2026-09-28

[E13] Jericho3-AI scheduled fabric

<https://investors.broadcom.com/news-releases/news-release-details/broadcom-unveils-industrys-highest-performance-fabric-ai>

Primary source; accessed by 2026-09-28