



AI Infrastructure Architecture Evolution

2015-2026

Silicon generations · Memory and interconnects
Conference disclosures · Systems and software

A generation-by-generation synthesis of conference disclosures and primary technical sources.
Evidence cutoff: 28 September 2026.
Complete English edition, paired with a full Simplified Chinese report.

Contents

01 Scope and executive synthesis	3
02 Reading the specifications	3
03 Axion: bringing the host into the platform	3
04 TPU v1 and the separate inference branch	5
05 TPU v2-v4: from chip to training computer	5
06 v5e, v5p and Trillium: two scaling objectives	5
07 Ironwood and the eighth-generation split	6
08 Accelerator generation comparison	7
09 ICI, Jupiter, Virgo and infrastructure offload	11
10 Software and integrated systems	15
11 Bandwidth balance and architectural trade-offs	21
12 Conference-to-product index	23
13 Naming crosswalk and remaining disclosure gaps	26
Source map	27
References	28

01 Scope and executive synthesis

Coverage: TPU and ICI are deep; Axion, offload and datacenter networking are full at the publicly disclosed level; video acceleration is proportional. Period: 2015-28 September 2026. The latest architecture disclosures reviewed include Cloud Next 2026, the 2026 training-TPU retrospective, and the public Hot Chips 2026 program. Pixel Tensor and edge Coral are excluded. Availability refers to internal deployment or cloud service availability, explicitly distinguished below.

Google's central architectural asset is continuity of the compiler-controlled tensor machine inside a changing system. The progression is from a latency-constrained inference coprocessor, to a distributed training computer, to workload-specific training and serving systems. Axion and Titanium extend that control to the host and infrastructure path. The important comparison with a GPU vendor is therefore the complete path from graph compilation to job placement, memory movement, collective communication and recovery, rather than matrix throughput alone.

02 Reading the specifications

All compute figures are per physical TPU package and dense unless stated otherwise; TFLOPS counts a multiply-add as two operations. Cloud names such as v3-8 count TensorCores, not eight physical chips. GB and GiB retain source units. HBM bandwidth is not bidirectional ICI bandwidth. Full pod footprint, usable allocation and maximum single-job slice can differ. A source's first deployment date is not automatically public cloud GA. n/d means not established in the reviewed primary documents; n/a means not applicable.

03 Axion: bringing the host into the platform

The first Axion C4A service became generally available in October 2024 and uses Neoverse V2. N4A, generally available on 27 January 2026, uses the N3 core and targets flexible general-purpose VM shapes. These are distinct workload positions, not proof that every N-series core supersedes every C-series core. Published VM limits describe sold resources; they do not establish physical die core counts, memory channels or socket topology.

Architectural implication: a host can bottleneck AI without executing the main matrix workload. Tokenization, input preparation, dispatch and checkpoint movement consume CPU and I/O resources; infrastructure processing competes with those tasks unless offloaded. Titanium removes portions of networking and storage work from the application CPU. The eighth-generation TPU disclosure explicitly includes Axion hosts. This creates a co-design opportunity, but it does not establish coherent CPU-TPU shared physical memory or uniform access latency.

Host CPU comparison · 1/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Codename	GA	Status	Core µarch	Cores / threads
Axion / C4A	n/d	2024-10	Shipping / available	Neoverse V2	n/d physical; VM ≤72 vCPU
Axion / N4A	n/d	2026-01-27	Shipping / available	Neoverse N3	n/d physical; VM ≤64 vCPU

Host CPU comparison · 2/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Process per die	Dies / packaging	L2 per core	L3 / domain
Axion / C4A	n/d	n/d	n/d	n/d
Axion / N4A	n/d	n/d	n/d	n/d

Host CPU comparison · 3/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Memory: channels; rate; GB/s	Sockets / links	PCIe / CXL	TDP	Vector / matrix ISA
Axion / C4A	DDR5; channels/rate n/d	n/d	n/d	n/d	Arm64; SVE2
Axion / N4A	DDR5; VM ≤512 GB; peak n/d	n/d	n/d	n/d	Arm64

04 TPU v1 and the separate inference branch

TPU v1 entered internal service in 2015. Its integer systolic array, software-managed buffers and deterministic execution addressed inference under strict response-time limits. The ISCA 2017 paper is valuable because it connects the machine to production workload measurements, including utilization and memory constraints. A large arithmetic array cannot compensate for insufficient weight reuse; the 34 GB/s off-chip interface made that limitation concrete. Its DDR-based memory and lack of a training pod distinguish it sharply from later HBM systems.

TPU v4i is a distinct inference design, deployed in 2020 and documented at ISCA 2021. Its 7 nm implementation and enlarged on-chip storage reflect an effort to retain useful working sets while controlling serving power. It should not be folded into the training TPU v4 row. More generally, the i, e and p branches represent different optimization objectives; a later number is insufficient to infer which device is best for a particular latency, capacity or scaling requirement.

05 TPU v2-v4: from chip to training computer

TPU v2 introduced BF16 training, HBM and a two-dimensional ICI torus. TPU v3 expanded matrix resources and memory capacity and moved to liquid cooling. Hot Chips 2020 explains the architectural organization; the later training retrospective connects it to larger systems and resilience. The important implementation consequence is that more compute in the same generation of manufacturing did not remain a chip-only change: power delivery, cooling, network size and the compiler's partitioning strategy had to evolve together.

TPU v4 combines stronger compute with SparseCore acceleration for embeddings and an optically reconfigurable interconnect. Optical circuit switches change physical connectivity at a coarse timescale; they are not packet switches executing each collective. Reconfiguration helps assemble healthy slices, vary topology and accommodate unavailable components. The 4,096-chip system is therefore also an availability architecture: at this scale, recovering usable capacity matters alongside reducing collective time. The ISCA 2023 paper supplies the explicit product link.

06 v5e, v5p and Trillium: two scaling objectives

TPU v5e is a cost-oriented training and inference platform with a smaller physical domain; v5p targets much larger training systems and more HBM per chip. Trillium (v6e) advances the efficiency branch with more compute and memory bandwidth. The e branch is not merely a partially disabled p chip: its domain size, memory budget and deployment economics should be evaluated on their own. For a model that fits in a small slice, paying for a much larger physical scale-up capability may yield little benefit.

The design question is the bottleneck after partitioning. Dense training can reuse weights enough to benefit from a larger MXU budget, whereas embedding traffic, small batches and decode expose memory and communication. A fair comparison holds model, precision, batch, latency objective and usable chip count constant. Comparing a complete v5p pod with a small Trillium slice answers a procurement-capacity question, not a per-chip architectural efficiency question.

07 Ironwood and the eighth-generation split

Ironwood/TPU7x introduces two compute chiplets with distinct local memory spaces. The package contains two TensorCores, four SparseCores and 192 GiB of HBM according to the cloud specification table. The compiler-visible locality differs from the earlier MegaCore abstraction. A larger package peak does not erase cross-chiplet movement: sharding and kernel placement now have an additional physical boundary to respect. The cloud documentation lists JAX and PyTorch support and explicitly excludes TensorFlow on TPU7x.

The April 2026 TPU 8 disclosure makes workload specialization explicit. TPU 8t retains a large torus for training; TPU 8i uses Boardfly for a smaller-diameter serving domain and adds a collective acceleration engine. The published 8i footprint is 1,152 chips, while the topology description separately states up to 1,024 active chips. These numbers must remain separate. Both products are recorded here as announced; the reviewed architectural announcement does not establish general availability.

Architectural interpretation: the split is evidence that domain size and network diameter have become independent product variables. A training system can tolerate more hops when long collective transfers amortize startup cost. Autoregressive inference repeatedly pays synchronization latency and may prefer fewer hops even with a smaller domain. More SRAM helps only the portion of the working set that actually fits; it does not imply that an arbitrary long-context KV cache resides entirely on chip.

08 Accelerator generation comparison

TPU family · 1/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Architecture	GA / milestone	Status	Process	Dies / packaging
TPU v1	INT8 systolic	2015 Q2	Historical shipping	28 nm	1
TPU v2	TensorCore	2017 Q3	Historical shipping	16 nm	1
TPU v3	TensorCore	2018 Q4	Historical shipping	16 nm	1
TPU v4i	Inference	2020 Q1	Historical shipping	7 nm	1
TPU v4	TensorCore + SparseCore	2021 deployment	Historical shipping	7 nm	1
TPU v5e	Efficiency	2023	Shipping / available	n/d	n/d
TPU v5p	Performance	2023 preview; GA n/d	Shipping / available	n/d	1
TPU v6e	Trillium	2024	Shipping / available	n/d	n/d
TPU7x	Ironwood	2025	Shipping / available	n/d	2 compute chiplets
TPU 8t	Training	2026-04-22	Announced; GA n/d	n/d	n/d
TPU 8i	Serving / CAE	2026-04-22	Announced; GA n/d	n/d	compute + CAE chiplet

TPU family · 2/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Compute units	Dense BF16	Dense FP8	Dense FP6 / FP4	FP64 vector / matrix
TPU v1	1 × 256×256	n/a	n/a	n/a	n/d
TPU v2	2 TC	46 TFLOPS	n/a	n/a	n/d
TPU v3	2 TC	123 TFLOPS	n/a	n/a	n/d
TPU v4i	1 TC; 4 MXU	138 TFLOPS	n/a	n/a	n/d
TPU v4	2 TC; 4 SC	275 TFLOPS	n/a	n/a	n/d
TPU v5e	1 TC	197 TFLOPS	n/d	n/a	n/d
TPU v5p	2 TC; 4 SC	459 TFLOPS	459 TFLOPS	n/a	n/d
TPU v6e	1 TC; 2 SC	918 TFLOPS	918 TFLOPS	n/a	n/d
TPU7x	2 TC; 4 SC	2,307 TFLOPS	4,614 TFLOPS	n/a	n/d
TPU 8t	n/d	n/d	n/d	12.6 PFLOPS FP4; dense basis n/d	n/d
TPU 8i	2 TC; 1 CAE	n/d	n/d	10.1 PFLOPS FP4; dense basis n/d	n/d

TPU family · 3/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Memory: type; stacks; GB; TB/s	LLC / SRAM
TPU v1	DDR3; 8 GB; 0.034 TB/s	28 MB
TPU v2	HBM2; 16 GiB; 0.700 TB/s	32 MiB
TPU v3	HBM2; 4 stacks; 32 GiB; 0.900 TB/s	32 MiB
TPU v4i	HBM; 8 GB; 0.614 TB/s	144 MB
TPU v4	HBM2; 4 stacks; 32 GiB; 1.2 TB/s	32 MiB
TPU v5e	HBM; stacks n/d; 16 GB; 0.819 TB/s	n/d
TPU v5p	HBM2e; 6 stacks; 95 GiB usable; 2.765 TB/s	128 MiB
TPU v6e	HBM; stacks n/d; 32 GiB; 1.638 TB/s	n/d
TPU7x	HBM3e; 8 stacks; 192 GiB; 7.380 TB/s	128 MiB
TPU 8t	216 GB; 6.528 TB/s	128 MB
TPU 8i	288 GB; 8.601 TB/s	384 MB

TPU family · 4/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Scale-up: links; GB/s per direction; domain	Host link	Power / cooling	Form factor
TPU v1	n/a	PCIe; n/d	75 W chip TDP	Google platform
TPU v2	2D torus; 256 chips; direction n/d	PCIe; n/d	280 W chip TDP; air	Google platform
TPU v3	2D torus; 1,024 chips	PCIe; n/d	450 W chip TDP; liquid	Google platform
TPU v4i	2 × 400 Gb/s; direction n/d	PCIe; n/d	175 W chip TDP	Google platform
TPU v4	6 × 50 GB/s/direction; 4,096	PCIe; n/d	liquid; TDP n/d	Google platform
TPU v5e	256 chips; 400 GB/s bidirectional	PCIe; n/d	n/d	Google platform
TPU v5p	6 × 100 GB/s/direction; 8,960	PCIe; n/d	liquid; TDP n/d	Google platform
TPU v6e	256 chips; 800 GB/s bidirectional	PCIe; n/d	n/d	Google platform
TPU7x	6 × 100 GB/s/direction; 9,216	PCIe; n/d	liquid; TDP n/d	Google platform
TPU 8t	3D torus; 9,600	PCIe; n/d	liquid; TDP n/d	Google platform
TPU 8i	Boardfly: 1,152 physical / ≤1,024 active	PCIe; n/d	liquid; TDP n/d	Google platform

09 ICI, Jupiter, Virgo and infrastructure offload

ICI is the accelerator fabric; Jupiter is the datacenter network. The Jupiter Rising and Jupiter Evolving papers explain the progression from centrally controlled Clos fabrics to optical reconfiguration in the datacenter. These disclosures should not be conflated with the TPU's local ICI topology. Virgo adds a training-oriented scale-out fabric in the TPU 8 era. A million-chip distributed training claim describes aggregation across systems, not a million-chip coherent memory domain or a single ICI pod.

Titanium is an infrastructure offload system, not a publicly documented merchant DPU SKU with a complete per-chip datasheet. Falcon describes a reliable hardware transport; a transport specification is not itself a Google-branded NIC product. The appropriate report boundary is to map the function and deployment path, retaining n/d for undisclosed silicon implementation. Argos video acceleration belongs alongside this portfolio because it shows the same workload-driven specialization at warehouse scale, but its video engines are not TPU tensor cores.

Networking and offload roles · 1/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	GA / milestone	Status	Class
Titanium	n/d	Shipping / available	Infrastructure offload
Jupiter	n/d	Shipping / available	Datacenter fabric
Falcon	n/d	Shipping / available	Transport; not a NIC SKU
Virgo	n/d	Announced; GA n/d	AI scale-out fabric

Networking and offload roles · 2/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Ports × speed	SerDes	Host interface	Programmable engine
Titanium	n/d	n/d	n/d	n/d
Jupiter	n/d	n/d	n/d	n/d
Falcon	n/d	n/d	n/d	n/d
Virgo	n/d	n/d	n/d	n/d

Networking and offload roles · 3/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Embedded CPU	On-card memory	Transport / RDMA	Congestion control
Titanium	n/d	n/d	n/d	n/d
Jupiter	n/d	n/d	Ethernet	n/d
Falcon	n/d	n/d	RDMA	n/d
Virgo	n/d	n/d	n/d	n/d

Networking and offload roles · 4/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Key offloads	Power
Titanium	Networking/storage	n/d
Jupiter	Fabric control	n/d
Falcon	Reliable transport	n/d
Virgo	Training fabric	n/d

ICI evolution · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Interconnect	Year / generation	Lane rate	Lanes / link	GB/s per link per direction
ICI	v2	n/d	n/d	n/d
ICI	v3	n/d	n/d	n/d
ICI	v4	n/d	n/d	50
ICI	Ironwood	n/d	n/d	100
ICI / Boardfly	TPU8i	n/d	n/d	n/d

ICI evolution · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Interconnect	Links / device	Topology	Domain limit	Coherence / semantics
ICI	n/d	2D torus	256	Explicit collectives
ICI	n/d	2D torus	1024	Explicit collectives
ICI	6	3D torus / OCS	4096	Explicit collectives
ICI	6	3D torus	9216	Explicit collectives
ICI / Boardfly	n/d	Boardfly	1152 physical / ≤1024 active	Explicit collectives / CAE

10 Software and integrated systems


XLA preserves a relatively stable programming contract while hardware changes. JAX expresses distributed computation; Pallas and Mosaic expose lower-level control for kernels whose data movement dominates. SparseCore and collective engines have value only when the software places and schedules work for them. Therefore portability and peak utilization are separate milestones: executing an unchanged model does not mean that its partitioning, fusion and scratchpad schedule exploit a new generation.

Software evolution

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.


Release / layer	Date	Hardware	Capability
XLA / TensorFlow	2017-2020	v2-v3	Graph compilation and placement
JAX / XLA	2021-2026	v4 onward	Distributed array programming
Pallas / Mosaic	2025-2026	Ironwood / TPU8	Explicit kernel-level locality
Native PyTorch	2026-04	TPU8 stack	Preview at disclosure

2017-2020: the training pod




Logical roles, not a physical wiring diagram.

2021-2025: reconfigurable scale



OCS changes connectivity; computation remains in the TPU.

2026: separate training and serving systems



Announced architecture; public GA is not inferred from this diagram.

System domains · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Platform	Year / milestone	CPU / accelerator count	Scale-up domain
TPUv4 pod	2021	4096 TPU chips	4096
Ironwood pod	2025	9216 TPU chips	9216
TPU8t	2026 announcement	9600 TPU chips	9600
TPU8i	2026 announcement	1152 physical / ≤1024 active	Boardfly

System domains · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Platform	Scale-out per accelerator	Rack power	Cooling
TPUv4 pod	n/d	n/d	Liquid
Ironwood pod	n/d	n/d	Liquid
TPU8t	n/d	n/d	Liquid
TPU8i	n/d	n/d	Liquid

11 Bandwidth balance and architectural trade-offs

Derived ratios below use memory bandwidth divided by dense compute, in bytes/FLOP. Capacity divided by compute has units of byte-seconds/FLOP and is deliberately reported separately in the canonical data. v5p uses 95 GiB of usable cloud memory, while the retrospective describes 96 GiB physical capacity. Ironwood bandwidth is 7,380 GB/s in the current cloud table; the retrospective rounds it differently. These are source-basis differences, not additional hardware generations.

Derived balance metrics

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation	HBM BW / BF16	HBM BW / FP8	One-way ICI / HBM	Performance / W
v2	0.015217	n/a	n/d	n/d: comparable power basis
v3	0.007317	n/a	n/d	n/d: comparable power basis
v4	0.004364	n/a	0.250	n/d: comparable power basis
v5p	0.006024	0.006024	0.217	n/d: comparable power basis
v6e	0.001784	0.001784	0.244	n/d: comparable power basis
v7	0.003199	0.001599	0.081	n/d: comparable power basis

Architectural interpretation: from v2 to Ironwood, BF16 compute grows about 50 times while HBM bandwidth grows about 10.5 times. The required reuse rises accordingly. Larger arrays shift responsibility toward compiler tiling, fusion and local storage. At system scale, optical reconfiguration trades a more complex control plane for improved slice assembly and fault accommodation. At the host boundary, direct I/O reduces copies but makes end-to-end scheduling and access protection more consequential. The practical evaluation should therefore include sustained throughput during communication, fault recovery time and the cost of recompiling or retuning models, not only isolated GEMM.

Capacity and host balance

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Metric	Value / calculation	Basis
TPUv4 HBM capacity/BF16	$32 \times 2^{30} / (275 \times 10^{12}) = 1.249 \times 10^{-4} \text{ B}\cdot\text{s}/\text{FLOP}$	Per chip; GiB converted to bytes
Ironwood HBM capacity/BF16	$192 \times 2^{30} / (2307 \times 10^{12}) = 8.936 \times 10^{-5} \text{ B}\cdot\text{s}/\text{FLOP}$	Per package; dense BF16
Axion DDR/core; L3/core	n/d	Physical CPU specifications incomplete

12 Conference-to-product index

Verified disclosure map · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Year	Paper / talk / disclosure
ISCA	2017	In-Datcenter Performance Analysis of a Tensor Processing Unit
ISCA	2021	Ten Lessons From Three Generations Shaped Google’s TPUv4i
ISCA	2023	TPU v4: An Optically Reconfigurable Supercomputer
Hot Chips	2020	Google training chips TPUv2 and TPUv3
IEEE Micro	2026	Google’s Training Supercomputers from TPU v2 to Ironwood
SIGCOMM	2015 / 2022	Jupiter Rising / Jupiter Evolving
ASPLOS	2021	Warehouse-Scale Video Acceleration
Cloud Next	2026	TPU 8t / 8i architecture
Hot Chips	2026	The Eighth Generation TPU Family

Verified disclosure map · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Presenter / authors	Product	Evidence type	Source
ISCA	Jouppi et al.	TPU v1	Product architecture and measurements	C01
ISCA	Jouppi et al.	TPU v4i	Product implementation and serving	C02
ISCA	Jouppi et al.	TPU v4	Architecture / system	C03
Hot Chips	Norrie / Patil	v2-v3	Architecture	C04
IEEE Micro	Google authors	v2-v7	Implementation / system retrospective	C05
SIGCOMM	Google authors	Jupiter	Network systems	C06 C07
ASPLOS	Ranganathan et al.	Argos	Product co-design	C08
Cloud Next	Gupta / Mugazambi	TPU8	Product announcement	P08
Hot Chips	Google	TPU8	Program verified; not recording analysis	C09

ISCA, ASPLOS, Hot Chips and IEEE Micro provide direct product disclosures in the selected corpus. MICRO and HPCA research should be included only when a documented product connection exists; none is asserted here. An ISSCC/JSSC circuit-paper-to-TPU mapping was not established in this pass, so circuit parameters are drawn from the identified product papers rather than attributed to ISSCC. Cloud Next and OCP are useful for launches and systems; MWC, Computex, GTC and other cloud events do not automatically add independent evidence for Google-owned silicon.

ISSCC context and implementation topics · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Year	Paper / talk / disclosure
ISSCC plenary	2018	50 Years of Computer Architecture: From Mainframe CPUs to Neural-Network TPUs
ISSCC plenary	2020	The Deep Learning Revolution and Its Implications for Computer Architecture and Chip Design
ISSCC Forum 3	2026	Power Delivery Trends and Demands in Data-Center AI Processors

ISSCC context and implementation topics · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Presenter / authors	Product	Evidence type	Source
ISSCC plenary	David Patterson	TPU / domain specialization	Official plenary record	C10
ISSCC plenary	Jeff Dean	TPU / chip-design context	Author companion paper and official record	C11 C10
ISSCC Forum 3	Houle Gan	Google datacenter processors	Official program; not a per-generation circuit datasheet	C12

ISSCC contributes both the domain-specialization rationale and a power-delivery implementation perspective. The 2026 forum entry establishes Google’s participation in that engineering discussion, but does not by itself reveal a particular TPU voltage-regulator circuit. These records complement the ISCA product papers and Hot Chips architecture disclosures rather than replacing them with inferred circuit specifications.

13 Naming crosswalk and remaining disclosure gaps

Names and boundaries

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Name	Maps to	Boundary
Trillium	TPU v6e	Efficiency branch
Ironwood	TPU7x	Family name vs cloud configuration
Axion	C4A / N4A	CPU brand vs VM series
TPU v4i	Inference TPU	Not training TPU v4
TPU 8t / 8i	Training / serving	Separate architecture objectives

The largest unresolved implementation fields are the physical Axion package organization, exact TDP and per-die process information for recent TPUs, and complete link electrical parameters. These gaps limit package-level power and area comparisons but do not prevent analysis of the disclosed memory and network hierarchy. The next useful evidence would be a circuit-level implementation paper, a complete TPU 8 service specification, and measured software behavior across Ironwood’s chiplet boundary.

Source map

- 01 Scope and executive synthesis - C01, C03, C05, C09, P08, P10
- 02 Reading the specifications - C02, P07, P09
- 03 Axion: bringing the host into the platform - E01, P08, P10, P11
- 04 TPU v1 and the separate inference branch - C01, C02
- 05 TPU v2-v4: from chip to training computer - C03, C04, C05
- 06 v5e, v5p and Trillium: two scaling objectives - P04, P05, P06
- 07 Ironwood and the eighth-generation split - E03, P07, P08, P13
- 08 Accelerator generation comparison - C01, C02, C03, C04, C05, E02, E03, P04, P05, P06, P07, P08
- 09 ICI, Jupiter, Virgo and infrastructure offload - C03, C04, C05, C06, C07, C08, E05, P07, P08, P10, P14
- 10 Software and integrated systems - C03, C04, C05, P07, P08, P13
- 11 Bandwidth balance and architectural trade-offs - C03, C04, C05, P05, P07, P08, P10
- 12 Conference-to-product index - C01, C02, C03, C04, C05, C06, C07, C08, C09, C10, C11, C12, P08
- 13 Naming crosswalk and remaining disclosure gaps - C02, P06, P07, P08, P10

References

Product and technical documentation

[P04] Cloud TPU v5e

<https://docs.cloud.google.com/tpu/docs/v5e>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P05] Cloud TPU v5p

<https://docs.cloud.google.com/tpu/docs/v5p>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P06] Cloud TPU v6e (Trillium)

<https://docs.cloud.google.com/tpu/docs/v6e>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P07] TPU7x (Ironwood)

<https://docs.cloud.google.com/tpu/docs/tpu7x>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P08] Inside the eighth-generation TPU: An architecture deep dive

<https://cloud.google.com/blog/products/compute/tpu-8t-and-tpu-8i-technical-deep-dive>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P09] TPU architecture

<https://docs.cloud.google.com/tpu/docs/system-architecture-tpu-vm>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P10] Axion N4A generally available

<https://cloud.google.com/blog/products/compute/axion-based-n4a-vm-s-now-in-preview>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P11] Arm VMs on Compute

<https://docs.cloud.google.com/compute/docs/instances/arm-on-compute>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P13] Inside the Ironwood codesigned AI stack

<https://cloud.google.com/blog/products/compute/inside-the-ironwood-tpu-codesigned-ai-stack>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[P14] Falcon RDMA network profiles

<https://docs.cloud.google.com/vpc/docs/rdma-network-profiles>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

Conference records and implementation disclosures

[C01] In-Datcenter Performance Analysis of a Tensor Processing Unit, ISCA 2017

<https://arxiv.org/pdf/1704.04760>

Primary source; accessed by 2026-09-28

[C02] Ten Lessons From Three Generations Shaped Google's TPuv4i, ISCA 2021

https://markgottscho.com/cv/papers/2021_NJouppi_ISCA.pdf

Primary source; accessed by 2026-09-28

[C03] TPU v4, ISCA 2023

<https://arxiv.org/pdf/2304.01433>

Primary source; accessed by 2026-09-28

[C04] Google training chips TPuv2 and TPuv3, Hot Chips 2020

https://www.hc32.hotchips.org/assets/program/conference/day2/HotChips2020_ML_Training_Google_Norrie_Patil.v01.pdf

Primary source; accessed by 2026-09-28

[C05] Google's Training Supercomputers from TPU v2 to Ironwood, IEEE Micro 2026

<https://arxiv.org/pdf/2606.15870>

Primary source; accessed by 2026-09-28

[C06] Jupiter Rising, SIGCOMM 2015

<https://research.google/pubs/jupiter-rising-a-decade-of-clos-topologies-and-centralized-control-in-googles-datacenter-network/>

Primary source; accessed by 2026-09-28

[C07] Jupiter Evolving, SIGCOMM 2022

<https://research.google/pubs/jupiter-evolving-transforming-googles-datacenter-network-via-optical-circuit-switches-and-software-defined-networking/>

Primary source; accessed by 2026-09-28

[C08] Warehouse-Scale Video Acceleration, ASPLOS 2021

<https://research.google/pubs/warehouse-scale-video-acceleration-co-design-and-deployment-in-the-wild/>

Primary source; accessed by 2026-09-28

[C09] Hot Chips 2026 conference program

<https://hc2026.hotchips.org/program/conference/>

Official indexed record or related announcement; full document download unavailable.

[C10] ISSCC plenary archive: Patterson2018 / Dean2020

<https://www.isscc.org/isscc-plenary-videos>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[C11] Dean ISSCC2020 companion paper

<https://arxiv.org/abs/1911.05289>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[C12] ISSCC2026 Forum3: Google power delivery

<https://submissions.mirasmart.com/ISSCC2026/PDF/ISSCC2026AdvanceProgram.pdf>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

Launches, ecosystem and deployment

[E01] First Axion C4A generally available

<https://cloud.google.com/blog/products/compute/first-google-axion-processor-c4a-now-ga-with-titanium-ssd/>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E02] Ironwood and Axion portfolio, November 2025

<https://cloud.google.com/blog/products/compute/ironwood-tpus-and-new-axion-based-vms-for-your-ai-workloads>

Indexed record or searchable official page; product details cross-checked with the associated technical document.

[E03] Two chips for the agentic era, Cloud Next 2026

<https://blog.google/innovation-and-ai/infrastructure-and-cloud/google-cloud/eighth-generation-tpu-agentic-era/>

Primary source; accessed by 2026-09-28

[E05] Data center and global networks built for the AI era

<https://cloud.google.com/blog/products/networking/data-center-and-global-networks-built-for-ai-era>

Indexed record or searchable official page; product details cross-checked with the associated technical document.