

Marvell

AI Infrastructure Architecture Evolution

2015-2026

Silicon generations · Memory and interconnects
Conference disclosures · Systems and software

A generation-by-generation synthesis of conference disclosures and primary technical sources.
Evidence cutoff: 28 September 2026.
Complete English edition, paired with a full Simplified Chinese report.

Contents

01 Scope and executive synthesis	3
02 Scope and numerical conventions	3
03 ThunderX: merchant CPUs to custom compute	3
04 OCTEON: infrastructure execution and acceleration	5
05 Teralynx: acquired switching architecture at AI scale	5
06 Optical DSPs, CPO and Photonic Fabric	5
07 Structera, Alaska P and the memory hierarchy	6
08 Fixed-schema product comparisons	7
09 Integrated systems and software	12
10 Derived balance and engineering implications	16
11 Conference index and product crosswalk	17
Source map	19
References	20

01 Scope and executive synthesis

Deep coverage: OCTEON, Teralynx, optical connectivity and memory/scale-up interconnect. Historical ThunderX CPUs are covered through their move toward custom solutions. Custom XPU's are covered only where implementation capabilities or customer links are public. Period: 2015-28 September 2026. Celestial AI and XConn are included with their pre-acquisition lineage because both acquisitions closed in February 2026. Storage and security controllers are proportional; enterprise access switching and wireless RF are outside the deep scope.

Marvell is best understood as a supplier of the data-movement and implementation substrate surrounding compute. Its CPU history matters because processor and SoC capability reappears in DPUs, near-memory acceleration and custom silicon. Its optical and electrical portfolio spans distances from a package boundary to a datacenter interconnect. The architectural question is how these components change the cost of placing compute, memory and network endpoints apart, and which software layer can exploit that freedom.

02 Scope and numerical conventions

A DPU core count is not a server-CPU product, and a near-memory accelerator is not a general AI training GPU. Switch bandwidth, optical-module capacity, SerDes lane rate and CXL payload bandwidth describe different resources. Lane GT/s, coded Gb/s and useful GB/s are retained separately. Product announcements and sampling targets do not establish volume availability. n/d denotes an unestablished field and n/a a nonapplicable field. Demonstrations announced for a future event are not treated as completed deployments.

03 ThunderX: merchant CPUs to custom compute

ThunderX originated at Cavium, which joined Marvell in 2018. ThunderX2 established a dual-socket Arm server platform and reached production deployments including Microsoft's internal Azure development servers. ThunderX3 was disclosed as a 96-core, four-thread-per-core design with eight DDR4-3200 channels and PCIe 4. The subsequent strategy statement shifted emphasis to hyperscaler-specific solutions. The 2020 ThunderX3 disclosure is therefore recorded as announced/repositioned, not proof of a continuing mass-market CPU roadmap.

Architectural interpretation: four-way SMT increases the number of runnable contexts without multiplying memory channels or execution units by four. Its benefit depends on latency hiding, instruction supply and interference. For a buy-side comparison, threads must not be compared directly with single-threaded Arm cores. The useful continuity is expertise in coherent fabrics, RAS and memory systems; it does not imply that current OCTEON cores are the same microarchitecture as ThunderX.

Historical host CPU lineage · 1/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Codename	GA	Status	Core µarch	Cores / threads
ThunderX	n/d	2015-era	Historical shipping	Custom Armv8	n/d
ThunderX2	n/d	2018-era	Historical shipping	Custom Armv8	n/d
ThunderX3	n/d	n/d	Announced / repositioned	Custom Armv8	96 / 384

Historical host CPU lineage · 2/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Process per die	Dies / packaging	L2 per core	L3 / domain
ThunderX	n/d	n/d	n/d	n/d
ThunderX2	n/d	n/d	n/d	n/d
ThunderX3	TSMC 7P	n/d	n/d	n/d

Historical host CPU lineage · 3/3

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Generation / reference SKU	Memory: channels; rate; GB/s	Sockets / links	PCIe / CXL	TDP	Vector / matrix ISA
ThunderX	DDR4; n/d	n/d	n/d	n/d	Armv8
ThunderX2	DDR4; n/d	2; links n/d	n/d	n/d	Armv8
ThunderX3	8 × DDR4-3200; 204.8 GB/s	1-2; links n/d	PCIe 4; 64 lanes	n/d	4 × 128-bit Neon

04 OCTEON: infrastructure execution and acceleration

OCTEON's historical MIPS and subsequent Arm families target packet processing, security and storage rather than generic host replacement. OCTEON TX and TX2 carry that infrastructure lineage into Arm-based SoCs. OCTEON10 moves to 5 nm and Neoverse N2, combining CPU execution with packet, crypto and machine-learning engines. The platform whitepaper describes a coherent cache and memory system around those engines, making it an integrated data-plane computer rather than a NIC with only fixed offloads.

The integrated ML engine serves inline or near-data inference, including infrastructure and networking workloads. Its presence does not make OCTEON a direct substitute for an HBM training accelerator. OCTEON10 Fusion adds radio-oriented acceleration and belongs to the 5G/vRAN branch; MWC disclosures help explain that branch but should not be allowed to dominate the AI datacenter narrative. SKU-specific CPU count, memory interface and port configuration must remain separate rather than assembled into an imaginary maximum device.

05 Teralynx: acquired switching architecture at AI scale

The Teralynx line came through Innovium and complements Marvell's other switching families. Teralynx7 provides a 12.8 Tb/s baseline. Teralynx10 moves to 51.2 Tb/s and entered production in 2024. The 2026 T100 announcement describes 102.4 Tb/s on a monolithic 3 nm die and a sampling target. Low latency, telemetry and congestion management matter because a training job exposes synchronized bursts rather than only independent cloud flows.

Architectural interpretation: a monolithic switch avoids some internal die-to-die crossings, but still faces a perimeter, SerDes and power-delivery budget. A faster switch can reduce fabric tiers only if its port configuration matches endpoint speed and required uplinks. The relevant comparison against Tomahawk is therefore a topology at a fixed traffic matrix, including small-packet rate, buffer pressure, congestion response and optical power, not a single aggregate Tb/s label.

06 Optical DSPs, CPO and Photonic Fabric

Marvell's optical DSP progression includes the 5 nm Nova 1.6T platform, the 3 nm Ara platform and later use-case-specific derivatives. The 2026 connectivity disclosures add 2 nm demonstrations and products such as Libra; the ECOC announcement is recorded as an announced demonstration at this cutoff. DSP specialization reflects the fact that reach, modulation, equalization and power budgets differ across intra-datacenter and longer-distance links. A transmit-only design changes the module partition rather than eliminating all signal-processing requirements.

Celestial AI's Photonic Fabric, disclosed at Hot Chips 2025 before the acquisition, targets optical connectivity closer to compute and memory. Marvell completed the acquisition on 2 February 2026. The acquisition announcement anticipates revenue contribution in later fiscal periods, so the technology should not be described as broadly deployed merely because the transaction closed. The slides provide an architectural direction and module concepts; they do not establish every proposed memory configuration as a shipping product.

Architectural interpretation: optics addresses distance and bandwidth density; it does not make remote memory equivalent to local HBM. Serialization, protocol traversal, endpoint buffering and software placement remain. An optical memory fabric needs a precise contract for addressability, consistency, latency, failure isolation and allocation. The power saved in electrical reach must be balanced against lasers, thermal control, packaging yield and repair strategy.

07 Structera, Alaska P and the memory hierarchy

Structera X expands memory; Structera A adds near-memory compute. The A 2504 disclosure combines sixteen Neoverse V2 cores with four DDR5-6400 channels and a CXL2.0/PCIe5 ×16 host interface. This distinction is fundamental: attached DRAM bandwidth can exceed the host link, while local compute can operate on data without sending every byte back to the CPU. Compression may reduce traffic for compressible data but is not a fixed bandwidth multiplier for all workloads.

XConn closed on 10 February 2026 and adds PCIe/CXL switching. Subsequent Structera S announcements include a 260-lane PCIe6 switch and a CXL switch with Q3 2026 sampling targets. Alaska P retimers support the electrical reach between endpoints and these switches. A retimer, a switch, a memory expander and a near-memory processor occupy four different positions in the path; combining their headline bandwidths would be meaningless.

08 Fixed-schema product comparisons

DPU and switch families · 1/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	GA / milestone	Status	Class
OCTEON TX2	2020-era	Historical shipping	DPU
OCTEON10	2021 announced	Shipping / available	DPU
Teralynx7	pre-2021	Historical shipping	Switch
Teralynx10	2024 production	Shipping / available	Switch
Teralynx T100	2026 sampling target	Announced; GA n/d	Switch

DPU and switch families · 2/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Ports × speed	SerDes	Host interface	Programmable engine
OCTEON TX2	≤200G datapath	n/d	n/d	Packet / crypto engines
OCTEON10	SKU-dependent	56G class; SKU-dependent	PCIe 5	Packet / crypto / ML
Teralynx7	12.8 Tb/s	n/d	n/a	Switch pipeline
Teralynx10	51.2 Tb/s	100G class	n/a	Switch pipeline
Teralynx T100	102.4 Tb/s	n/d	n/a	Switch pipeline

DPU and switch families · 3/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Embedded CPU	On-card memory	Transport / RDMA	Congestion control
OCTEON TX2	12-36 Armv8	n/d	Ethernet	n/d
OCTEON10	Neoverse N2	DDR5; SKU-dependent	Ethernet	n/d
Teralynx7	n/d	n/d	Ethernet	Telemetry
Teralynx10	n/d	n/d	Ethernet	Congestion-aware routing
Teralynx T100	n/d	n/d	Ethernet	AI congestion control

DPU and switch families · 4/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Key offloads	Power
OCTEON TX2	Security / storage	n/d
OCTEON10	Infrastructure / inline ML	n/d
Teralynx7	L2/L3	n/d
Teralynx10	L2/L3 / telemetry	n/d
Teralynx T100	BGA / CPC / CPO	n/d

Accelerator boundary · 1/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Architecture	GA / milestone	Status	Process	Dies / packaging
Structera A 2504	Near-memory Arm compute	n/d	Announced; GA n/d	5 nm	n/d
Customer custom XPU	Customer-specific	n/d	Per customer program	n/d	n/d

Accelerator boundary · 2/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Compute units	Dense BF16	Dense FP8	Dense FP6 / FP4	FP64 vector / matrix
Structera A 2504	16 Neoverse V2	n/d	n/d	n/d	n/d
Customer custom XPU	n/d	n/d	n/d	n/d	n/d

Accelerator boundary · 3/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Memory: type; stacks; GB; TB/s	LLC / SRAM
Structera A 2504	DDR5; 4 channels; 6400 MT/s; 204.8 GB/s	n/d
Customer custom XPU	n/d	n/d

Accelerator boundary · 4/4

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Product	Scale-up: links; GB/s per direction; domain	Host link	Power / cooling	Form factor
Structera A 2504	n/a	CXL2.0 / PCIe5 ×16	n/d	CXL device
Customer custom XPU	n/d	n/d	n/d	n/d

Connectivity products · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Interconnect	Year / generation	Lane rate	Lanes / link	GB/s per link per direction
Alaska P	2024-2026	PCIe6: 64 GT/s	16	n/d
Structera S PCIe 60260	2026	PCIe6: 64 GT/s	260 lanes / switch	n/d
Structera X / A	2024	PCIe5: 32 GT/s	16	~63 before protocol overhead
Photonic Fabric	2025-2026	n/d	n/d	n/d

Connectivity products · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.


Interconnect	Links / device	Topology	Domain limit	Coherence / semantics
Alaska P	n/d	Retimed point-to-point	n/a	Protocol-aware retimer
Structera S PCIe 60260	n/d	Switched PCIe	n/d	PCIe
Structera X / A	1	CXL attachment	n/d	CXL memory; nonuniform latency
Photonic Fabric	n/d	Optical fabric	n/d	Module-specific; not assumed coherent

09 Integrated systems and software

Marvell's custom-silicon offering brings compute, packaging and interface IP together, but a customer program is not a public merchant accelerator SKU. The NVIDIA collaboration establishes an integration direction for custom platforms without proving that every custom XPU uses NVLink. System diagrams therefore show functional interfaces, not an invented Marvell-only rack. Storage controllers and cryptographic engines remain important for data ingestion, persistence and trust, even when they do not execute model layers.


OCTEON uses infrastructure software and acceleration APIs, Teralynx uses a switch SDK and supported SAI/SONiC integration, and CXL products depend on host firmware and operating-system memory management. These are separate software contracts. An open interface helps integration but does not automatically solve placement, paging or fault containment. Near-memory compute additionally needs an offload mechanism with a measured break-even point between moving data and moving work.

2015-2020 integrated roles



The second era includes announced and roadmap components; it is not a shipping rack bill of materials.

2021-2026 integrated roles



The second era includes announced and roadmap components; it is not a shipping rack bill of materials.

System boundary · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Platform	Year / milestone	CPU / accelerator count	Scale-up domain
Partner/custom platforms	2015-2026	Customer-specific	n/d

System boundary · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Platform	Scale-out per accelerator	Rack power	Cooling
Partner/custom platforms	n/d	n/d	n/d

Software integration contracts

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Layer	Interface	Function	Boundary
Switch control	SDK / supported open interfaces	Forwarding and telemetry	ASIC-specific implementation
Infrastructure	Firmware / driver / acceleration APIs	Packet, storage and security offload	SKU and platform dependent
System integration	Host software and partner stack	Resource and traffic management	No single vendor-only AI framework

10 Derived balance and engineering implications

Derived bandwidth hierarchy

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Resource	Peak calculation	Consequence
ThunderX3 DRAM/core	$8 \times 3200 \times 8 / 1000 / 96 = 2.133 \text{ GB/s}$	Per core; not per SMT thread
Structera A local DRAM	$4 \times 6400 \times 8 / 1000 = 204.8 \text{ GB/s}$	Local memory-controller peak
PCIe5 ×16	~63 GB/s each direction	Coding-adjusted; packet overhead extra
Local DRAM / host-link ratio	$204.8 / 63 \approx 3.25$	Local compute can avoid host-link bottleneck
Custom-XPU HBM / FLOP / W	n/d	Customer specifications not public

The strongest cross-generation trend is separation of concerns: host compute, infrastructure compute, near-memory compute and transport each acquire specialized silicon. This can improve efficiency but adds integration boundaries and failure domains. Memory expansion should be evaluated as a tier with latency and bandwidth asymmetry; optical scale-up should be evaluated as a system with endpoint and control overhead. The decisive measurements are workload-level traffic avoided, tail latency, recovery behavior and power for the complete link, not isolated component peaks.

11 Conference index and product crosswalk

Disclosure map · 1/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Year	Paper / talk / disclosure
Hot Chips	2020	ThunderX3
Hot Chips	2025	Photonic Fabric Module
ISSCC / VLSI	2022-2025	SerDes disclosures indexed by vendor
OFC / company events	2023-2026	Optical DSP evolution
ECOC	2026	2nm optical demonstrations

Disclosure map · 2/2

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Venue	Presenter / authors	Product	Evidence type	Source
Hot Chips	Marvell	ThunderX3	Vendor confirmation of presentation	E09 E08
Hot Chips	Phil Winterbottom	Celestial AI	Architecture; pre-acquisition	C01
ISSCC / VLSI	Marvell	SerDes IP	Vendor index; not product-wide attribution	C02
OFC / company events	Marvell	Nova / Ara	Implementation platform / launch	E10 E14
ECOC	Marvell	Optical connectivity	Announcement; event completion not asserted	E15

No direct ISCA/MICRO/HPCA/ASPLOS product-paper chain is asserted for current Marvell DPUs and switches in this edition. The conference emphasis shifts appropriately toward Hot Chips, ISSCC/VLSI signaling and OFC/ECOC integration. A SerDes test-chip result must not be copied into a product table without a documented implementation link. The appendix preserves the vendor’s conference index as a research lead, not as a substitute for reading each circuit paper.

Names and roles

n/d: not established; n/a: not applicable. Precision, scope and direction are shown in the cells; sources are mapped at the end.

Name	Lineage	Role
ThunderX	Cavium → Marvell	Historical server CPU / custom transition
Teralynx	Innovium → Marvell	Datacenter switches
Photonic Fabric	Celestial AI → Marvell	Optical scale-up platform
Structera S	XConn integration	PCIe / CXL switching
Structera A / X	Marvell	Near-memory compute / expansion

The material gaps are customer-XPU compute specifications, matched DPU power/SKU tables, confirmed volume availability of the newest switches, and fully measured optical-memory access behavior. Public information is sufficient to reconstruct portfolio direction, but not to rank a custom Marvell accelerator against a GPU on dense BF16 or performance per watt. That comparison requires the named customer system and its actual software configuration.

Source map

- 01 Scope and executive synthesis - E04, E06, E07, E09, E11, P02
- 02 Scope and numerical conventions - E11, E12, E15, P01, P03
- 03 ThunderX: merchant CPUs to custom compute - E08, E09, E17, P01, P05
- 04 OCTEON: infrastructure execution and acceleration - E01, P01, P02
- 05 Teralynx: acquired switching architecture at AI scale - E02, E03, P03
- 06 Optical DSPs, CPO and Photonic Fabric - C01, E06, E10, E14, E15
- 07 Structera, Alaska P and the memory hierarchy - E07, E11, E12, E13, E16, P06
- 08 Fixed-schema product comparisons - C01, E01, E02, E03, E04, E05, E06, E11, E12, E16, P01, P02, P03
- 09 Integrated systems and software - E03, E04, E05, E11, P02
- 10 Derived balance and engineering implications - C01, E04, E08, E11
- 11 Conference index and product crosswalk - C01, C02, E03, E04, E05, E06, E07, E08, E09, E10, E11, E12, E13, E14, E15, P03

References

Product and technical documentation

[P01] OCTEON10 DPU platform whitepaper

<https://www.marvell.com/content/dam/marvell/en/public-collateral/embedded-processors/marvell-octeon-10-dpu-platform-white-paper.pdf>

Primary source; accessed by 2026-09-28

[P02] Data processing units

<https://www.marvell.com/products/data-processing-units.html>

Primary source; accessed by 2026-09-28

[P03] Teralynx switching portfolio

<https://www.marvell.com/products/data-center-switches.html>

Primary source; accessed by 2026-09-28

[P05] Product selector guide July2020

https://www.marvell.com/content/dam/marvell/en/psg/marvell_psg.pdf

Primary source; accessed by 2026-09-28

[P06] Structera CXL products

<https://www.marvell.com/products/cxl.html>

Primary source; accessed by 2026-09-28

Conference records and implementation disclosures

[C01] Celestial AI Photonic Fabric Module, Hot Chips2025

https://hc2025.hotchips.org/assets/program/conference/day2/76-Hotchips_25_Celestial_AI_v4.pdf

Primary source; accessed by 2026-09-28

[C02] Custom AI investor event2025; SerDes conference index

<https://www.marvell.com/content/dam/marvell/en/company/assets/marvell-custom-ai-investor-event-2025.pdf>

Primary source; accessed by 2026-09-28

Launches, ecosystem and deployment

[E01] OCTEON10 introduction, 2021

<https://www.marvell.com/company/newsroom/marvell-extends-octeon-leadership-industry-first-5nm-dpu.html>

Primary source; accessed by 2026-09-28

[E02] Teralynx10 51.2T production

<https://www.marvell.com/company/newsroom/marvell-teralynx-512t-ethernet-switch-enters-volume-production-for-global-ai-cloud-deployments.html>

Primary source; accessed by 2026-09-28

[E03] Teralynx T100 102.4T

<https://investor.marvell.com/news-events/press-releases/detail/1024/marvell-announces-availability-of-industrys-first-102-4-tbps-switch-purpose-built-for-ai-and-cloud-data-center-infrastructure>

Primary source; accessed by 2026-09-28

[E04] Accelerated infrastructure for the AI era, 2025

<https://www.marvell.com/content/dam/marvell/en/company/assets/marvell-accelerated-infrastructure-for-the-ai-era-event.pdf>

Primary source; accessed by 2026-09-28

[E05] Custom AI infrastructure with NVIDIA

<https://investor.marvell.com/news-events/press-releases/detail/97/marvell-and-nvidia-to-provide-custom-solutions-for-advanced-ai-infrastructure>

Primary source; accessed by 2026-09-28

[E06] Celestial AI acquisition completion

<https://www.marvell.com/company/newsroom/marvell-completes-acquisition-of-celestial-ai.html>

Primary source; accessed by 2026-09-28

[E07] XConn acquisition completion

<https://investor.marvell.com/news-events/press-releases/detail/1007/marvell-completes-acquisition-of-xconn-technologies>

Primary source; accessed by 2026-09-28

[E08] ThunderX3 architecture

<https://www.marvell.com/blogs/the-next-generation-of-thunderx-delivers-performance-and-power-advantages-to-cloud-and-hpc-server-markets.html>

Primary source; accessed by 2026-09-28

[E09] ThunderX strategy change

<https://www.marvell.com/blogs/arm-processors-in-the-data-center.html>

Primary source; accessed by 2026-09-28

[E10] Teralynx10 and Nova platform

<https://investor.marvell.com/news-events/press-releases/detail/188/marvell-announces-cloud-optimized-51-2-tbps-networking-platform-for-ai/ml-and-data-center-networks>

Primary source; accessed by 2026-09-28

[E11] Structera A and X launch

<https://www.marvell.com/company/newsroom/marvell-introduces-breakthrough-structera-cxl-product-line-to-address-server-memory-bandwidth-and-capacity-challenges-in-cloud-data-centers.html>

Primary source; accessed by 2026-09-28

[E12] Structera S PCIe6 260-lane switch

<https://www.marvell.com/company/newsroom/marvell-260-lane-pcie-6-switch-ai-data-center-scale-up.html>

Primary source; accessed by 2026-09-28

[E13] Structera S CXL switch

<https://investor.marvell.com/news-events/press-releases/detail/1017/marvell-launches-next-generation-cxl-switch-enabling-memory-pooling-to-break-through-the-ai-memory-wall>

Primary source; accessed by 2026-09-28

[E14] Optical DSP platform evolution

<https://www.marvell.com/company/newsroom/marvell-1-6t-optical-dsp-ai-data-center-connectivity.html>

Primary source; accessed by 2026-09-28

[E15] ECOC2026 demonstration announcement

<https://www.marvell.com/company/newsroom/marvell-industry-first-2nm-optical-technology-ai-data-center-infrastructure-ecoc-2026.html>

Primary source; accessed by 2026-09-28

[E16] Alaska P adoption

<https://investor.marvell.com/news-events/press-releases/detail/1001/marvell-announces-adoption-of-its-pcie-retimers-by-leading-ai-and-data-center-infrastructure-providers>

Primary source; accessed by 2026-09-28

[E17] ThunderX2 deployment

<https://www.marvell.com/company/newsroom/marvells-thunderx2-solution-now-deployed-for-microsoft-azure-development.html>

Primary source; accessed by 2026-09-28